

Managing Arrays

with the OMR garbage collector

Prerequisites

THIS LAB IS “BRING YOUR OWN LAPTOP”

- You need:
 - Linux, osx, or windows laptop
 - a C++11 toolchain: msvc, clang, gcc
 - git
 - cmake
- Clone the skeleton project
`git clone --recursive https://github.com/rwy0717/splash2018-omr-gc`

Don't forget your laptop charger!

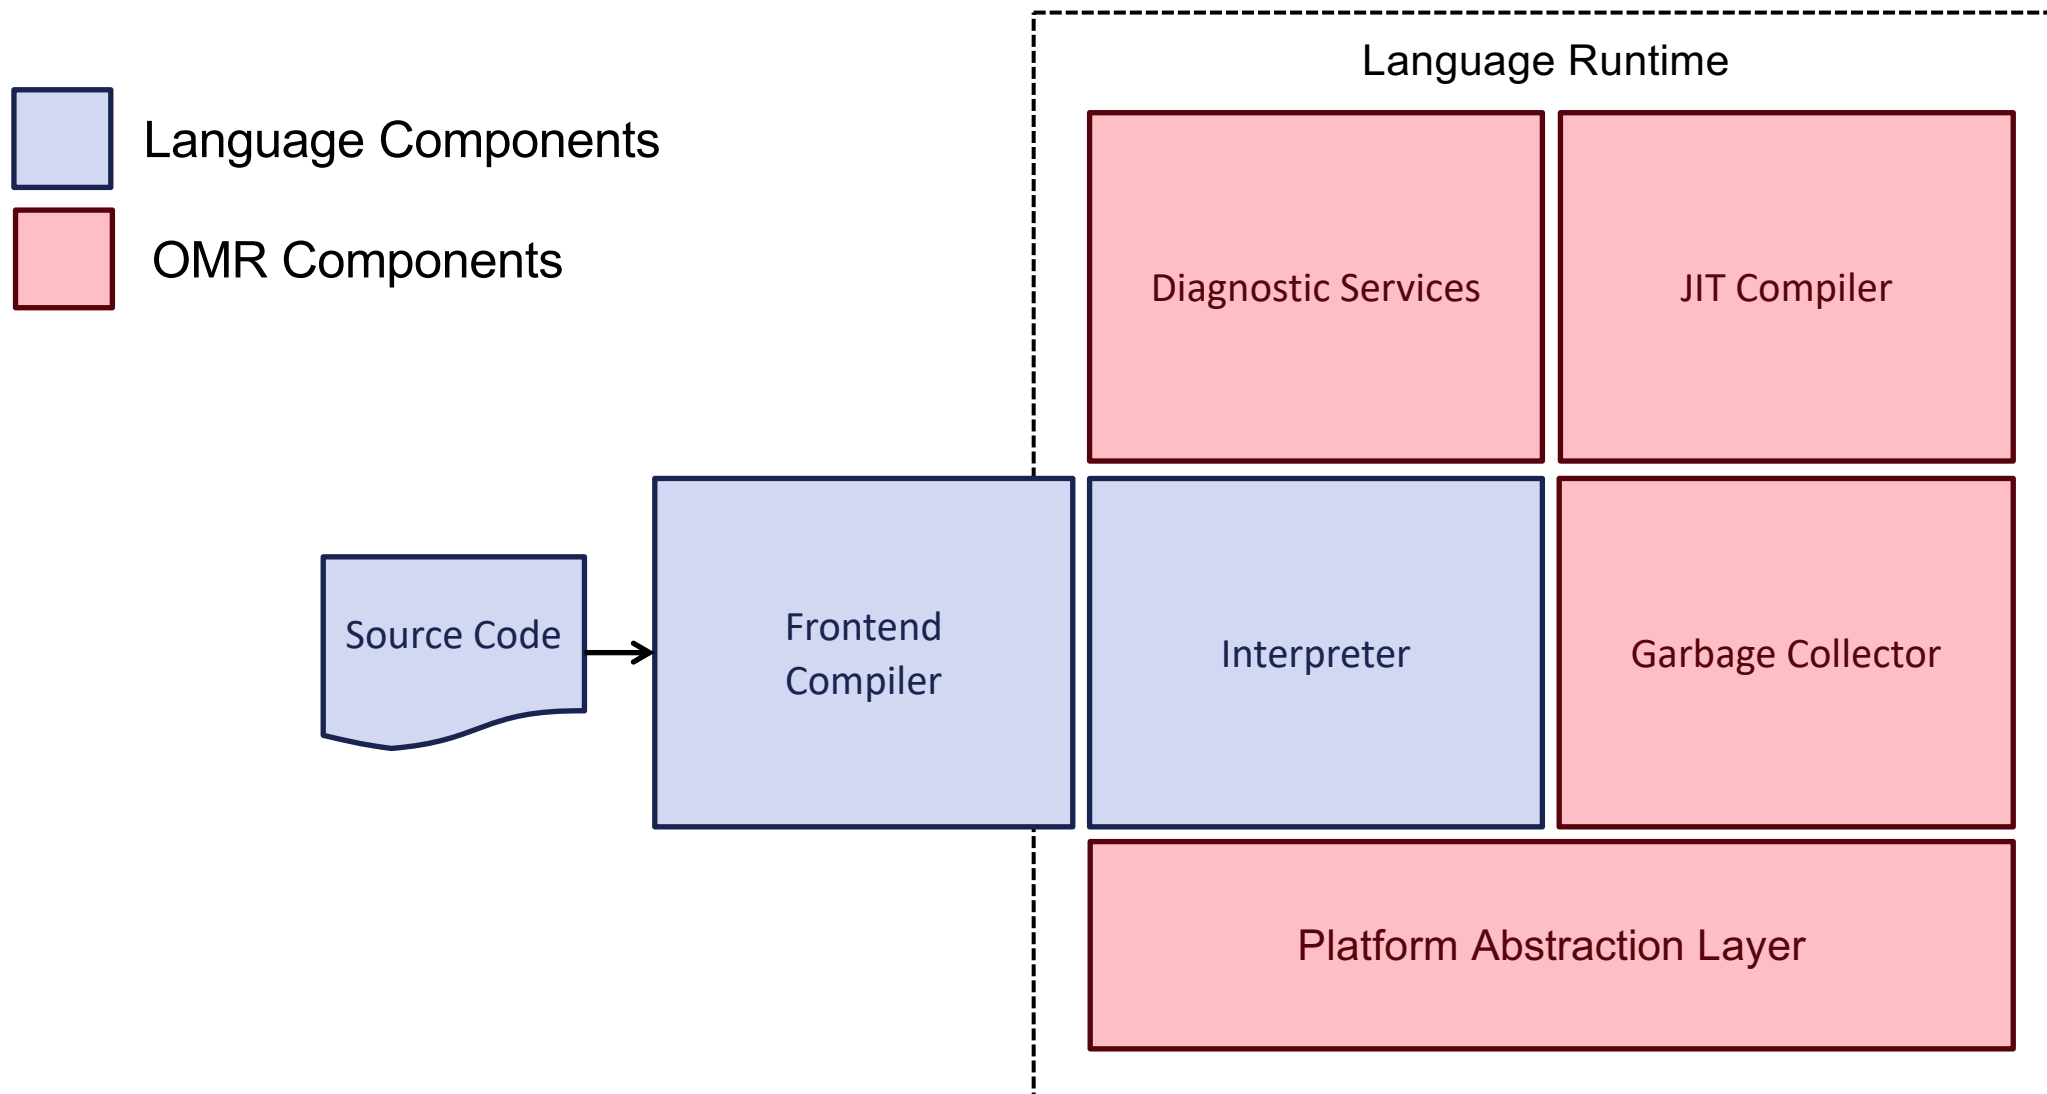
Today, we're going to implement a garbage collector for simple fixed-length arrays.

Eclipse OMR

A toolkit for building language runtimes

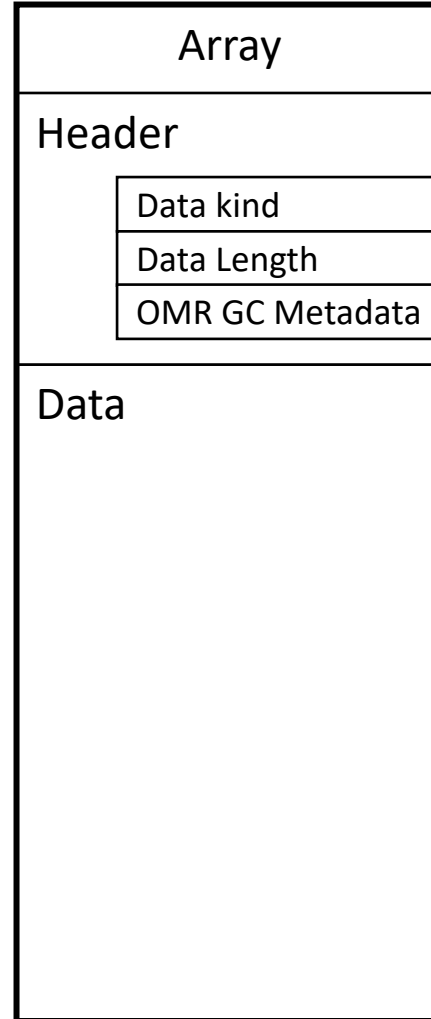
<https://github.com/eclipse/omr>



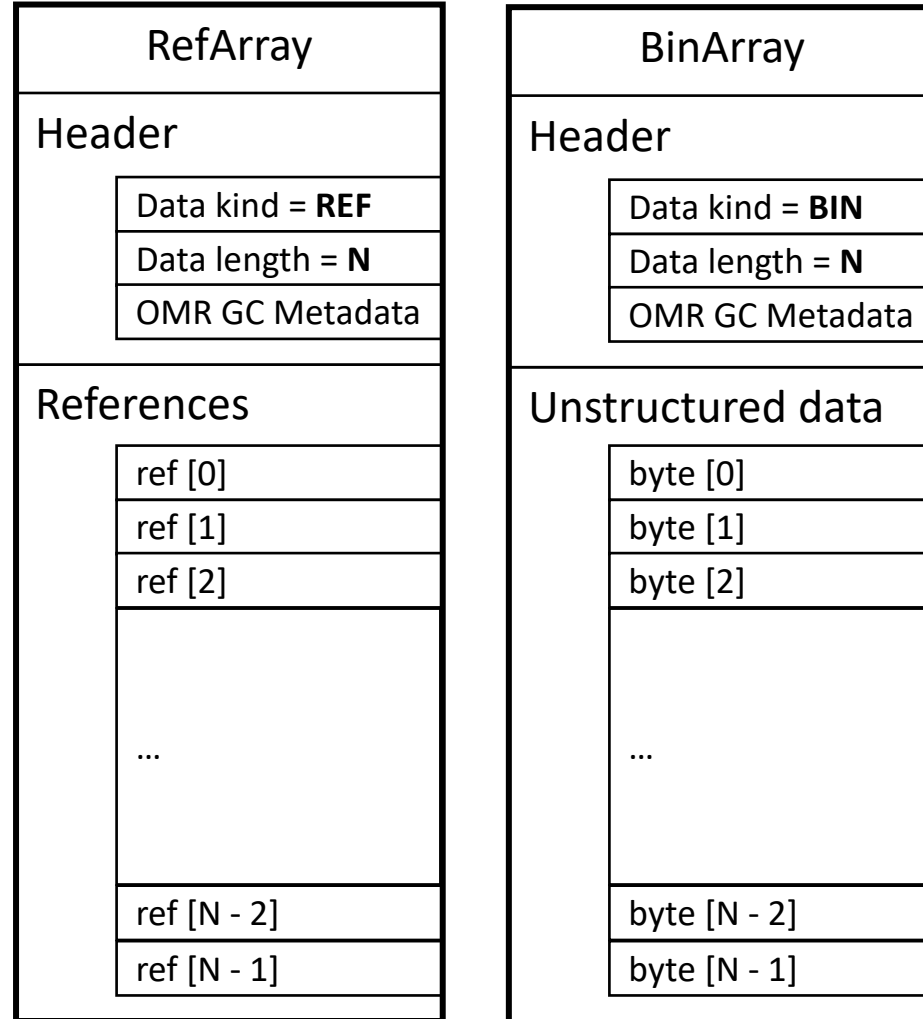


Our array object model

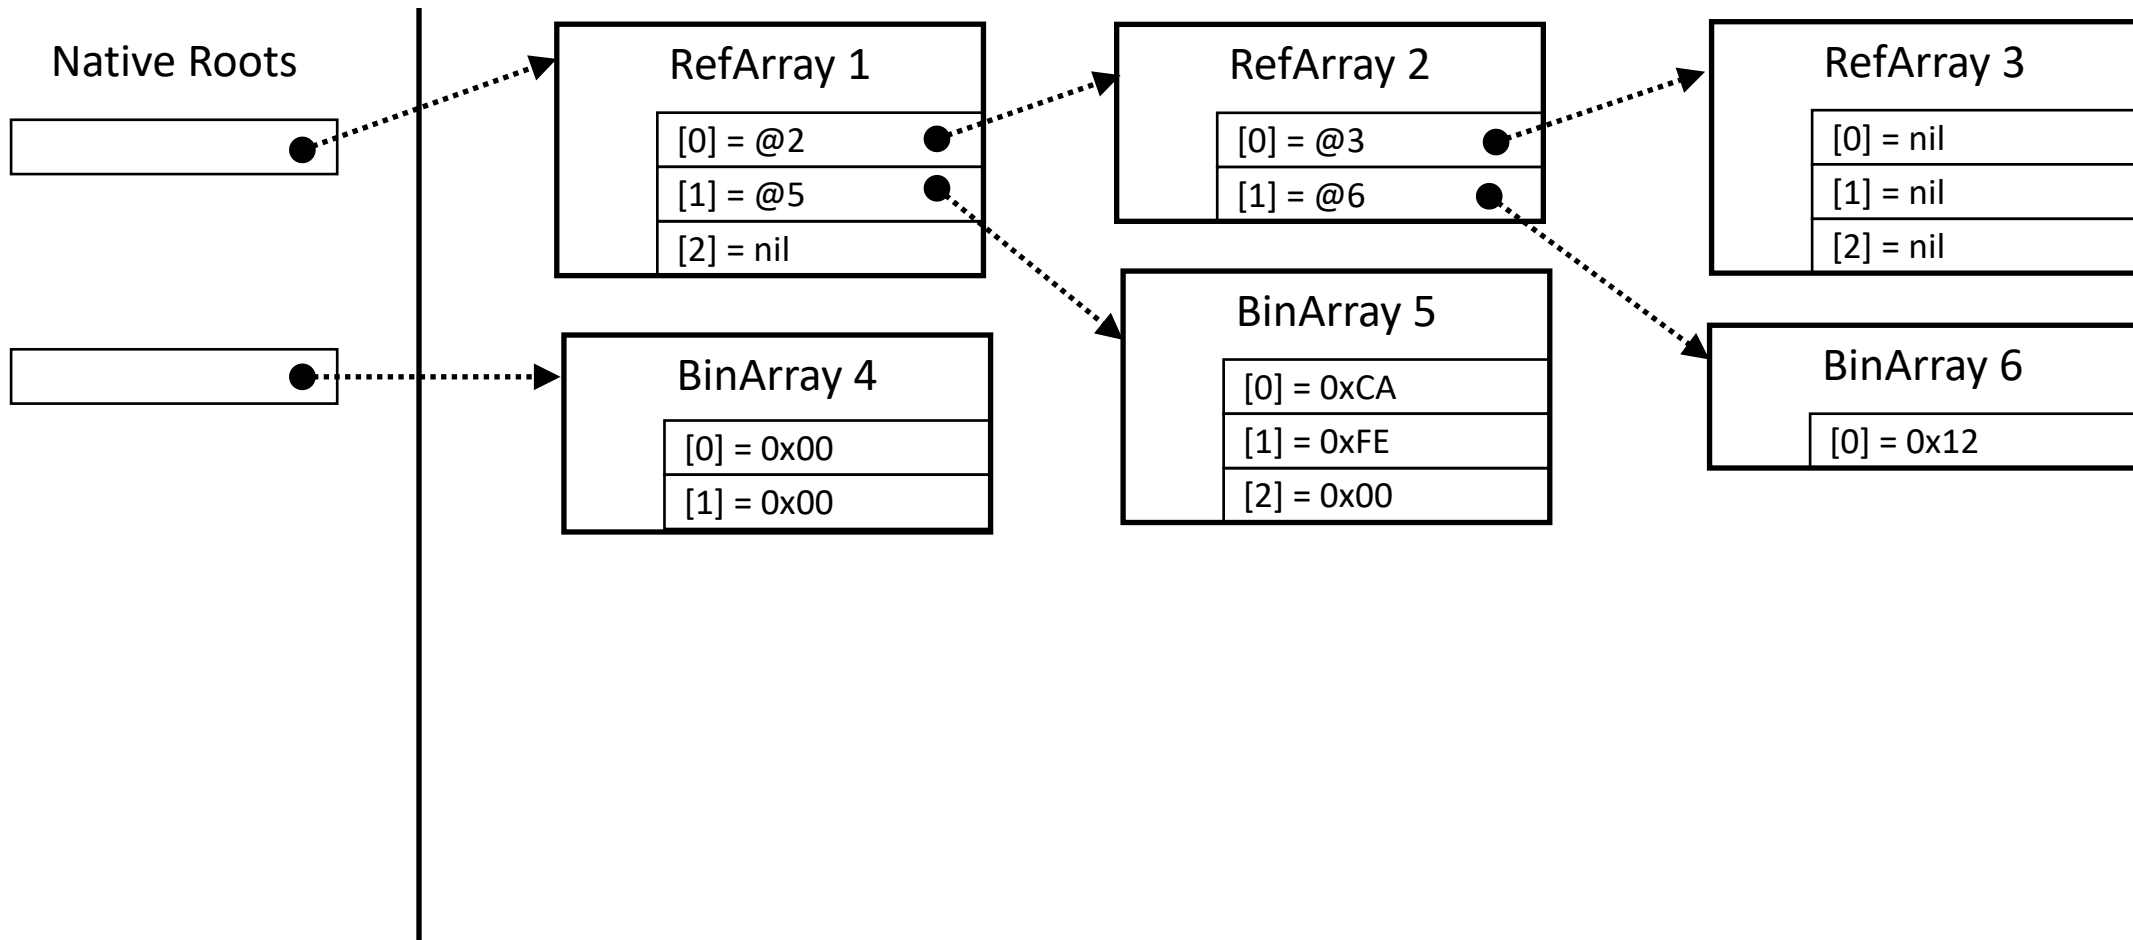
Our array object model:



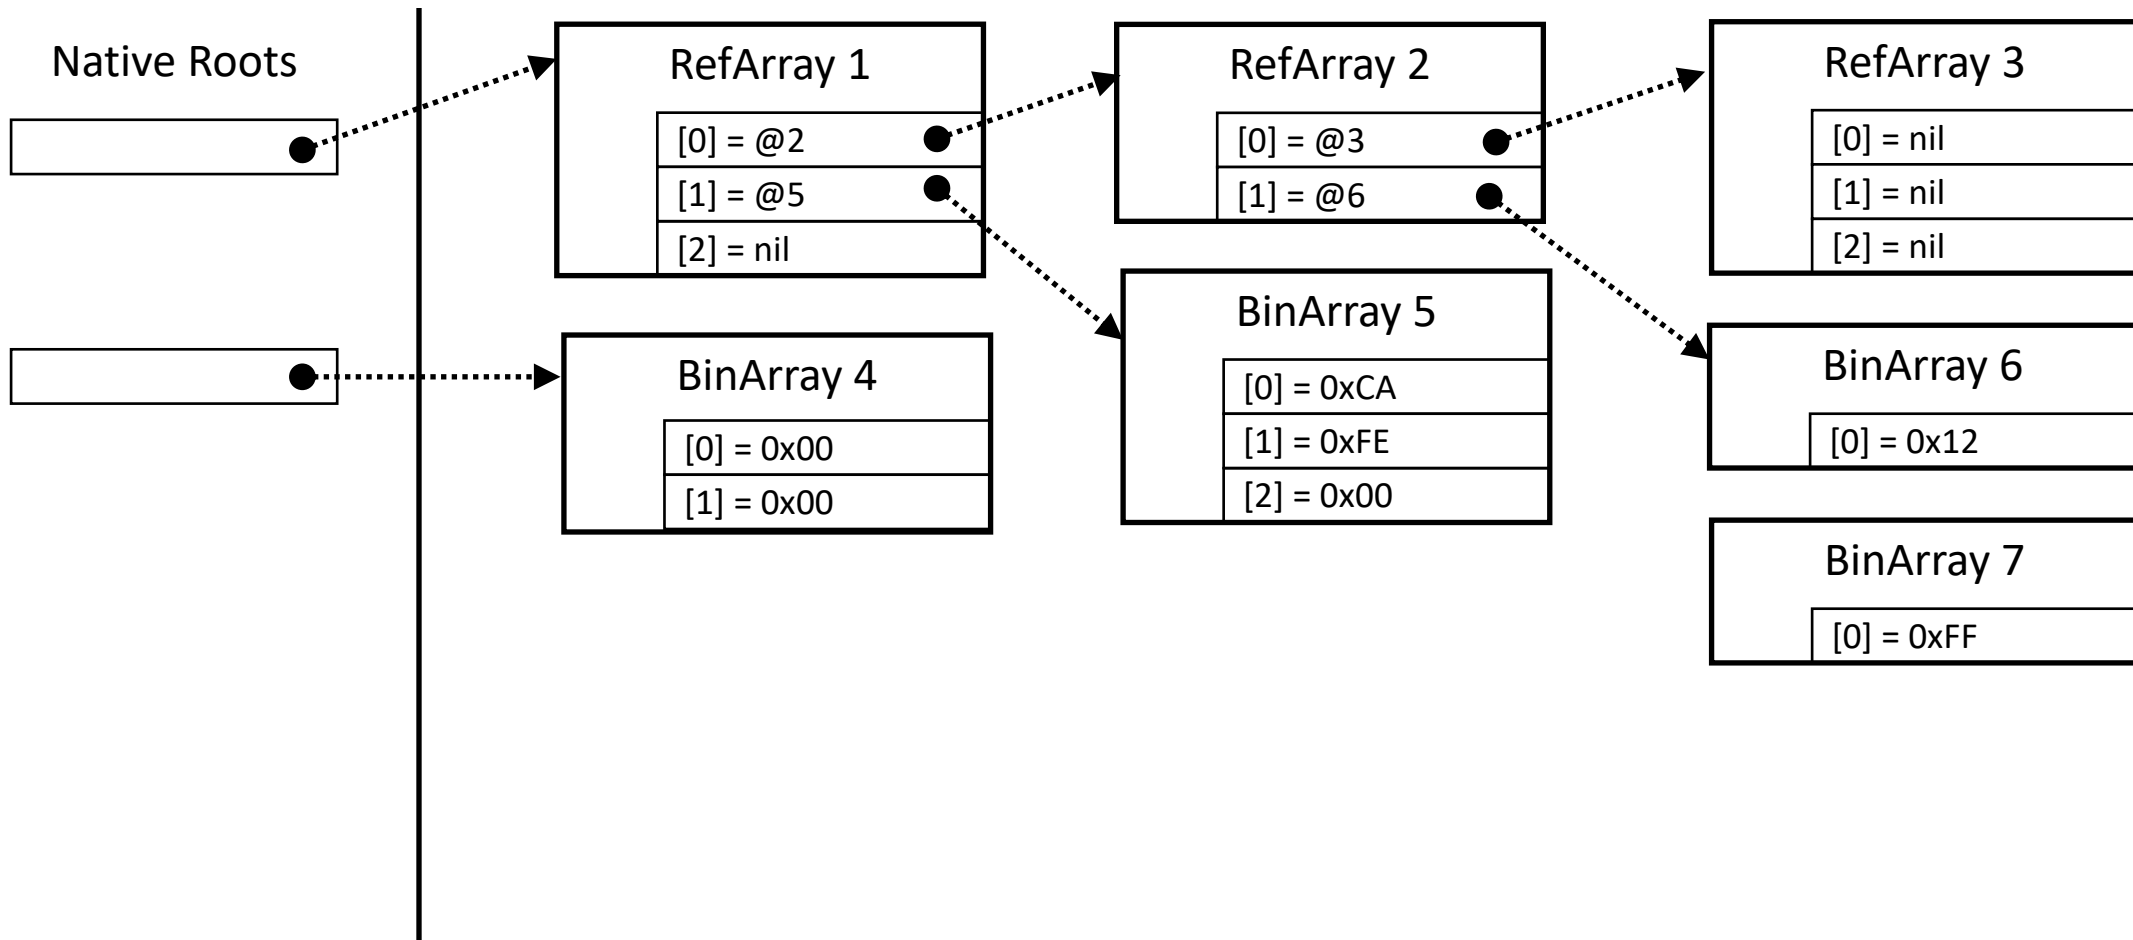
Two kinds of Data: References and bytes



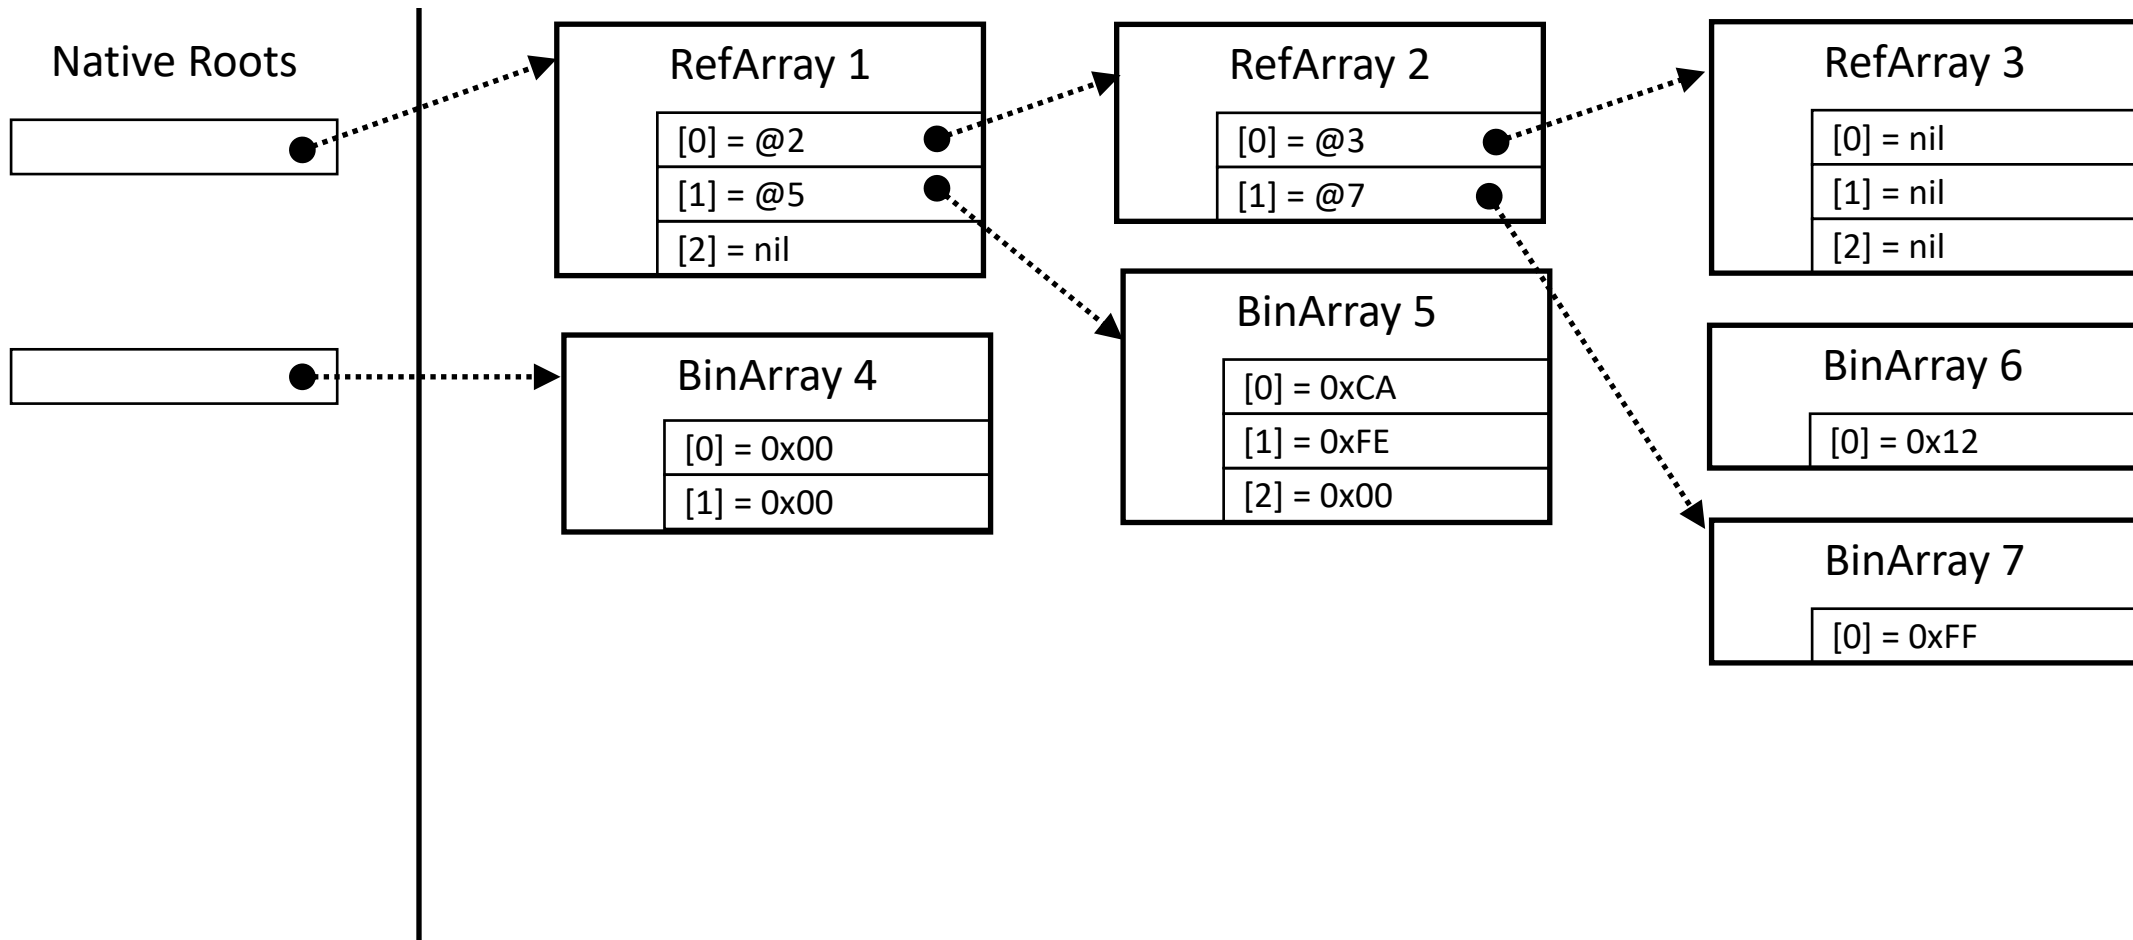
Memory: Graphs of objects



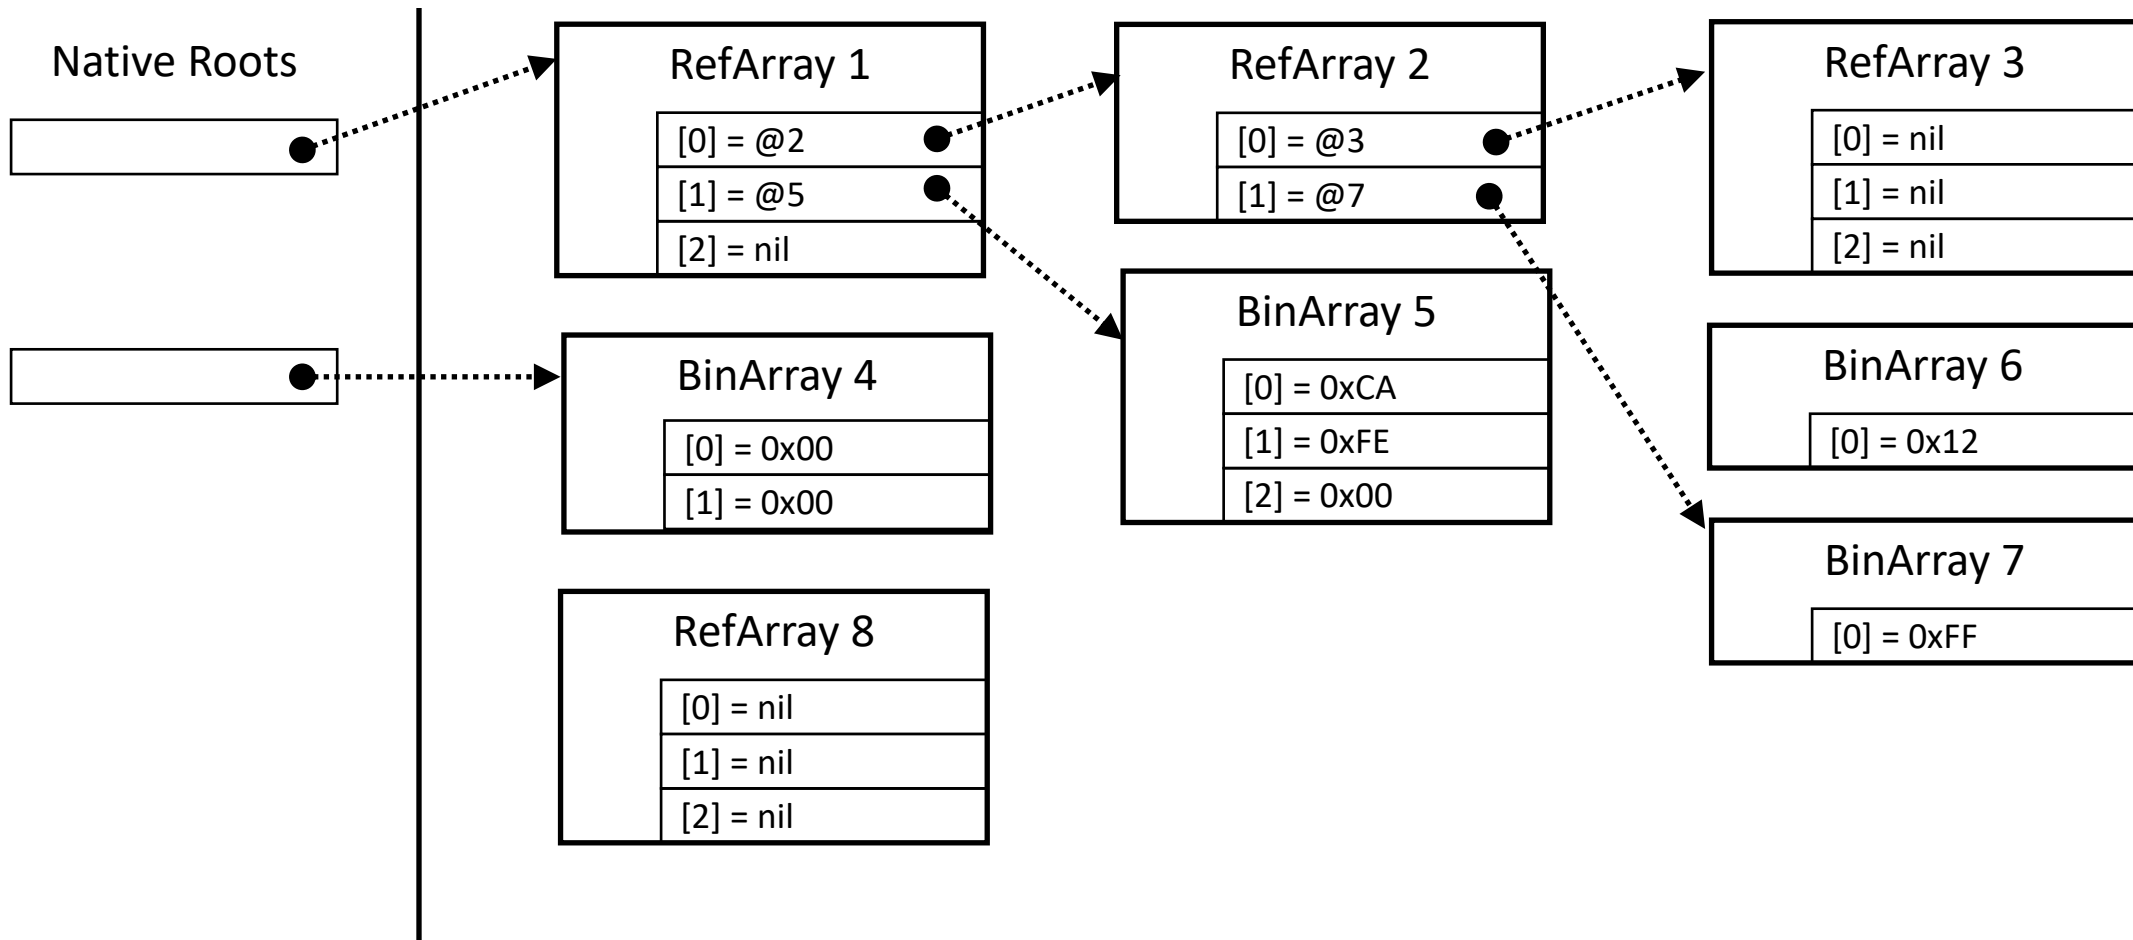
Memory: Graphs of objects



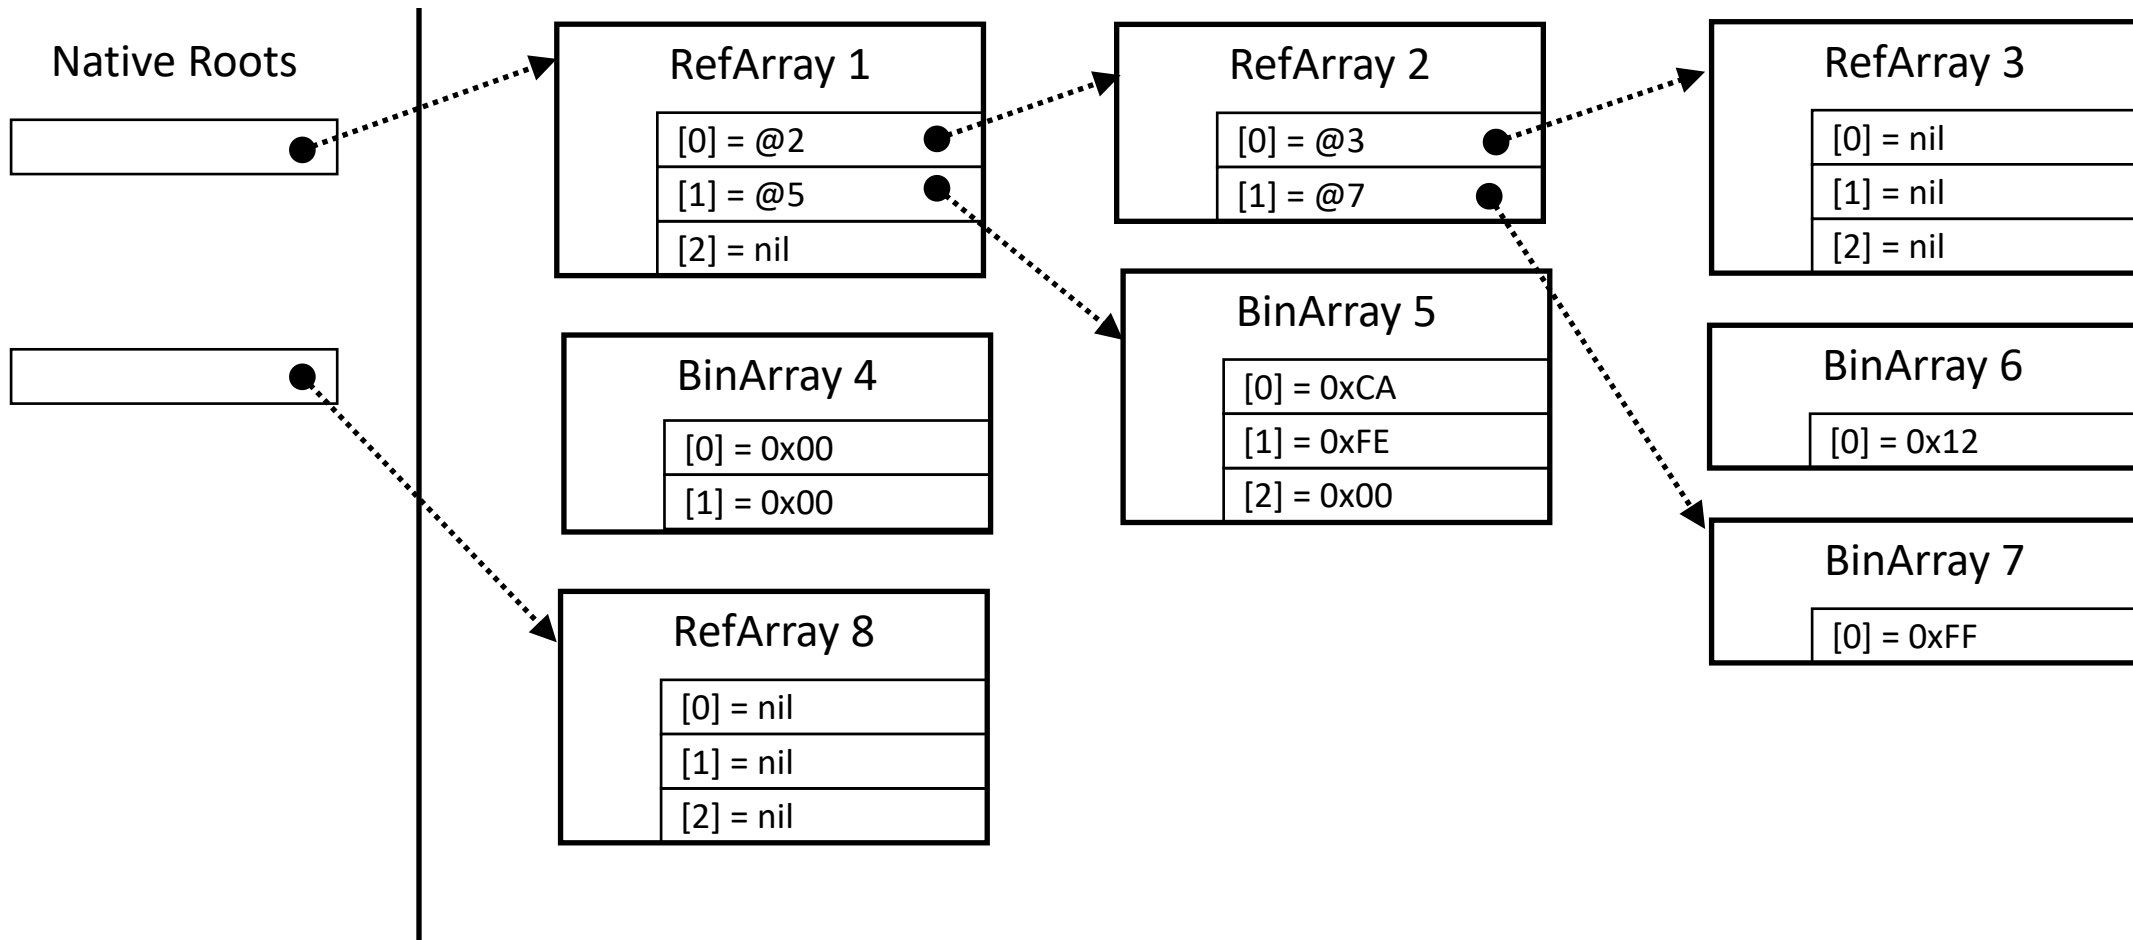
Memory: Graphs of objects



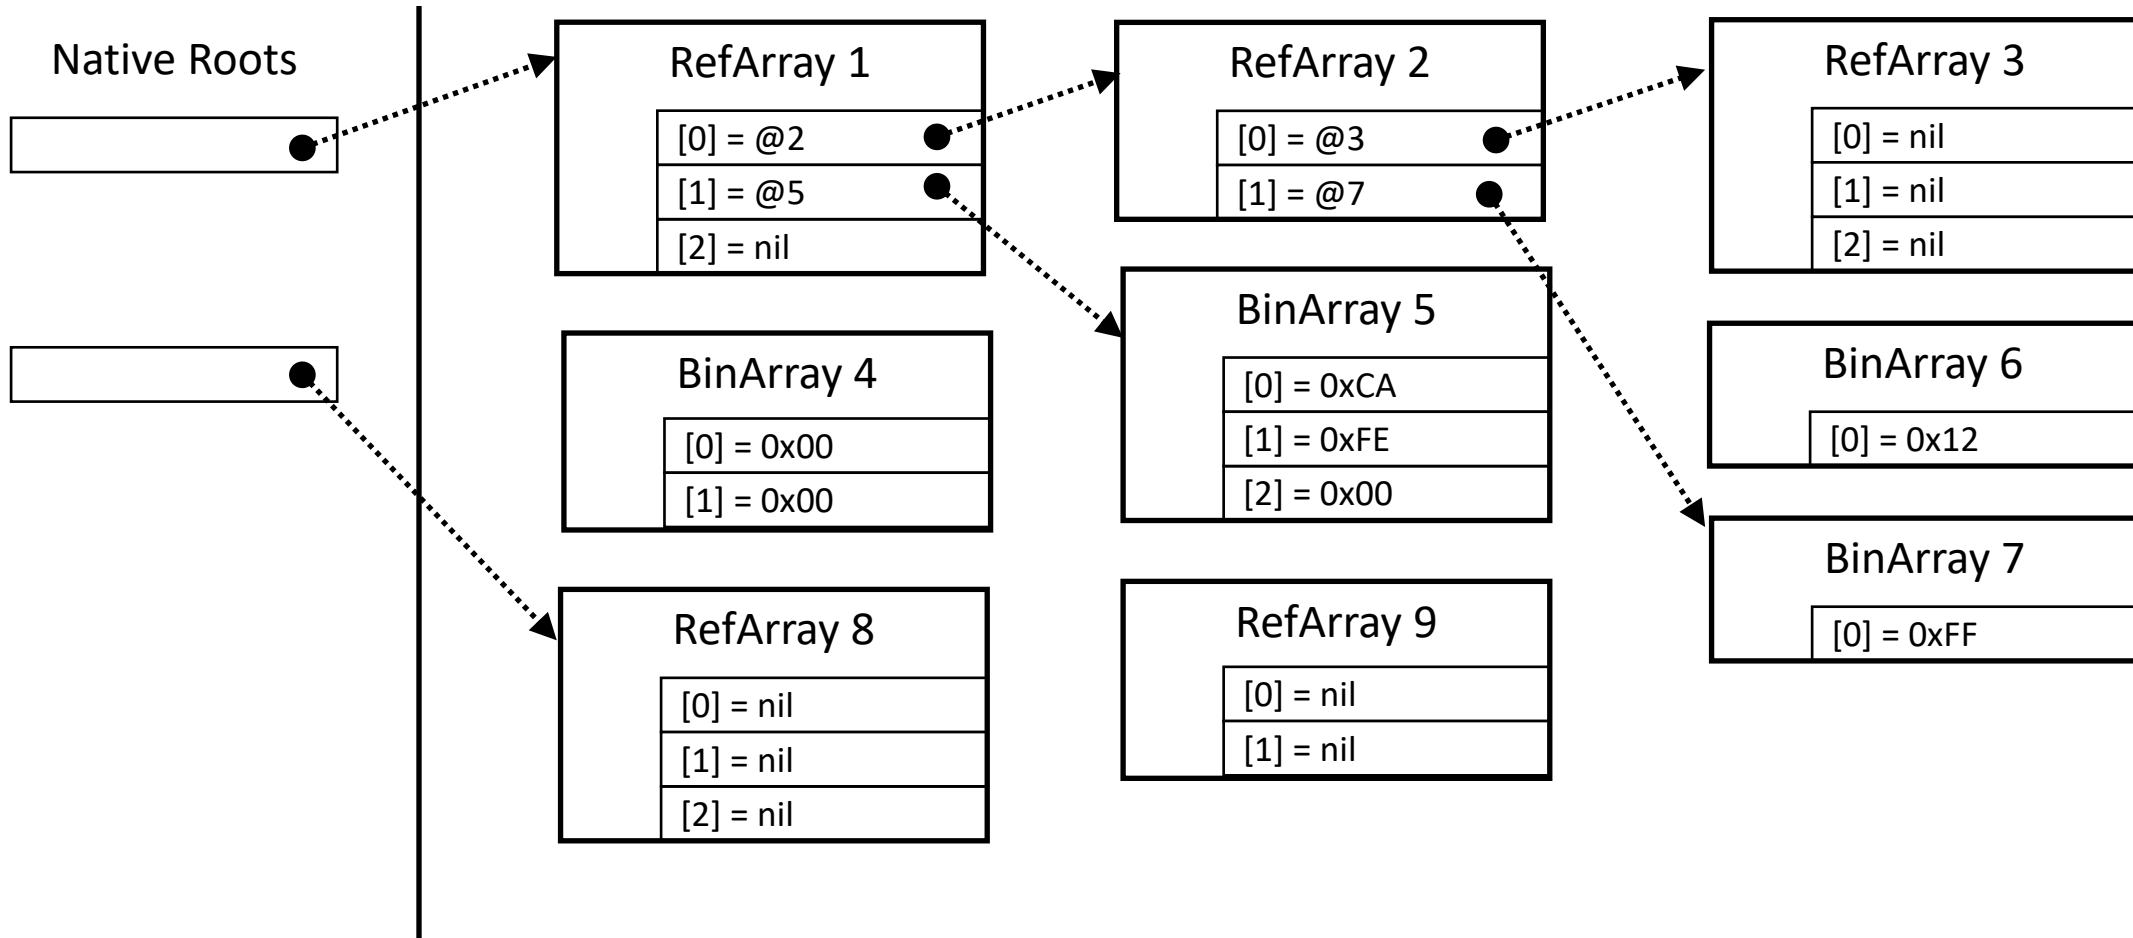
Memory: Graphs of objects



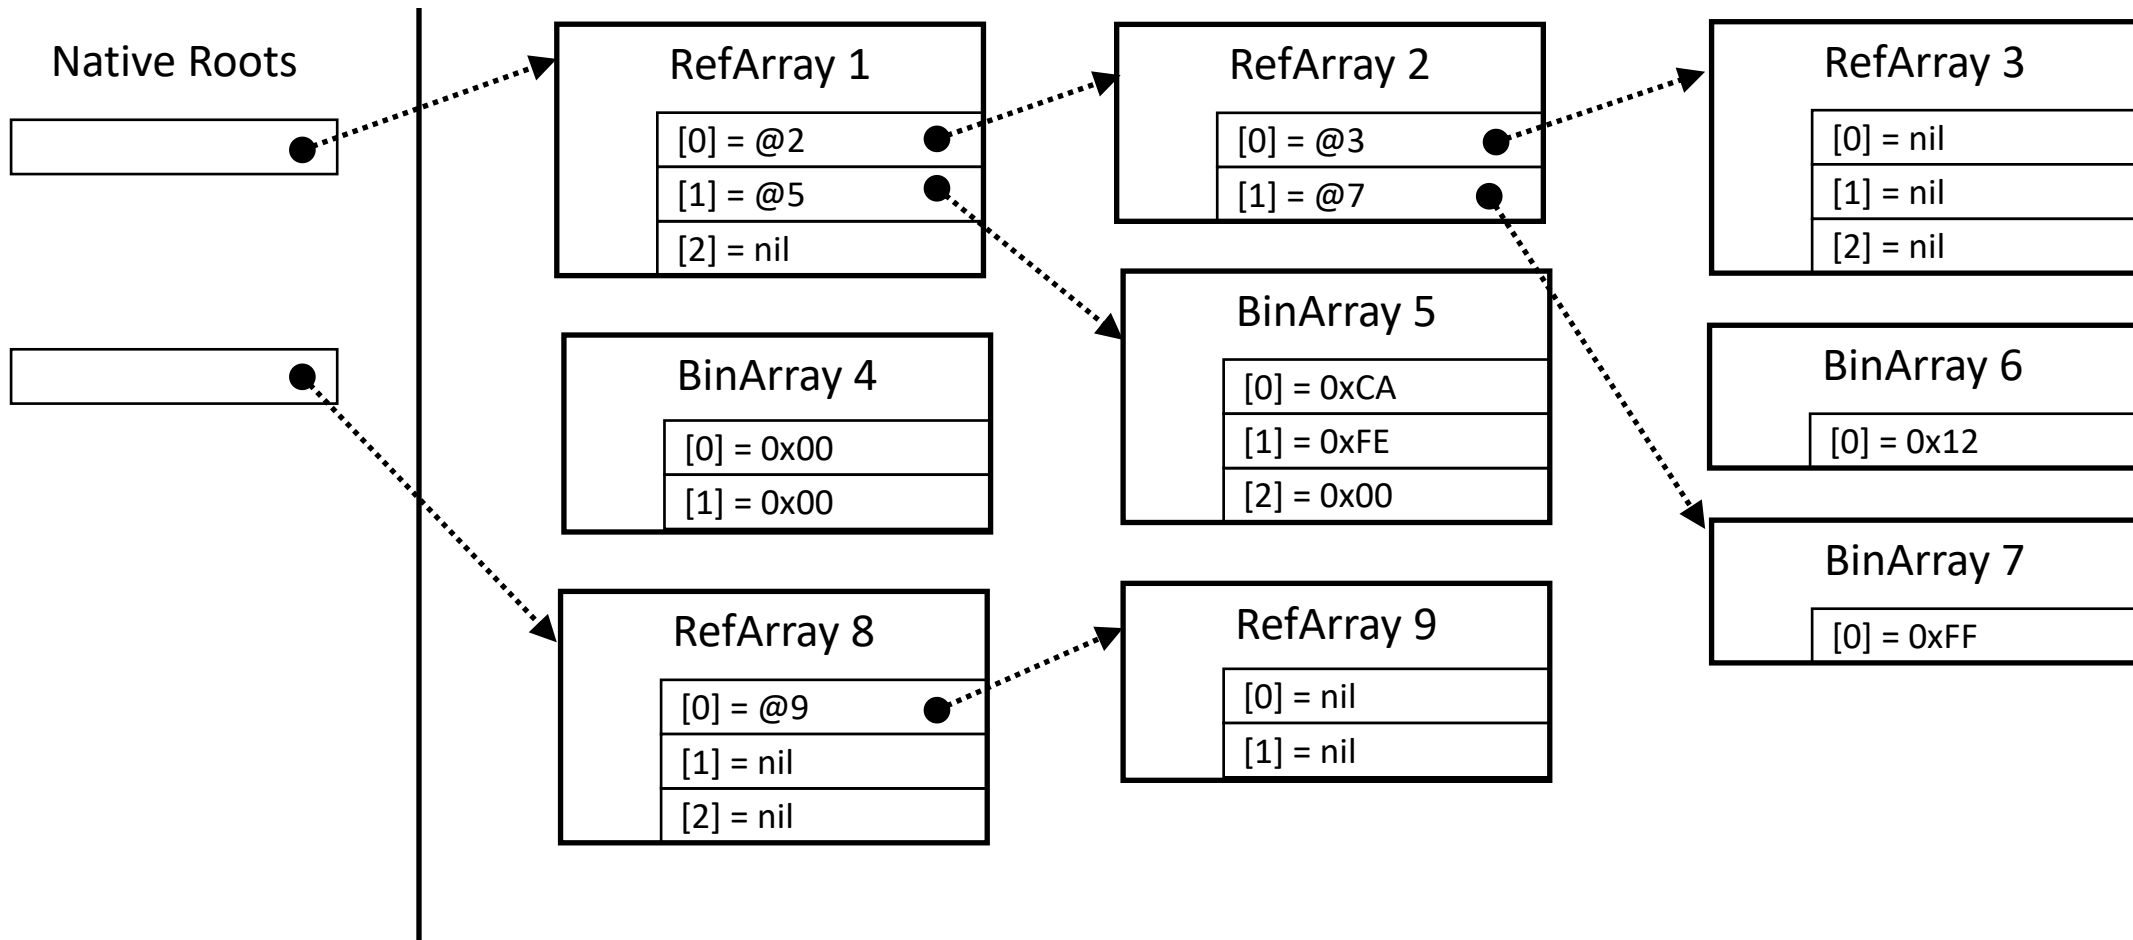
Memory: Graphs of objects



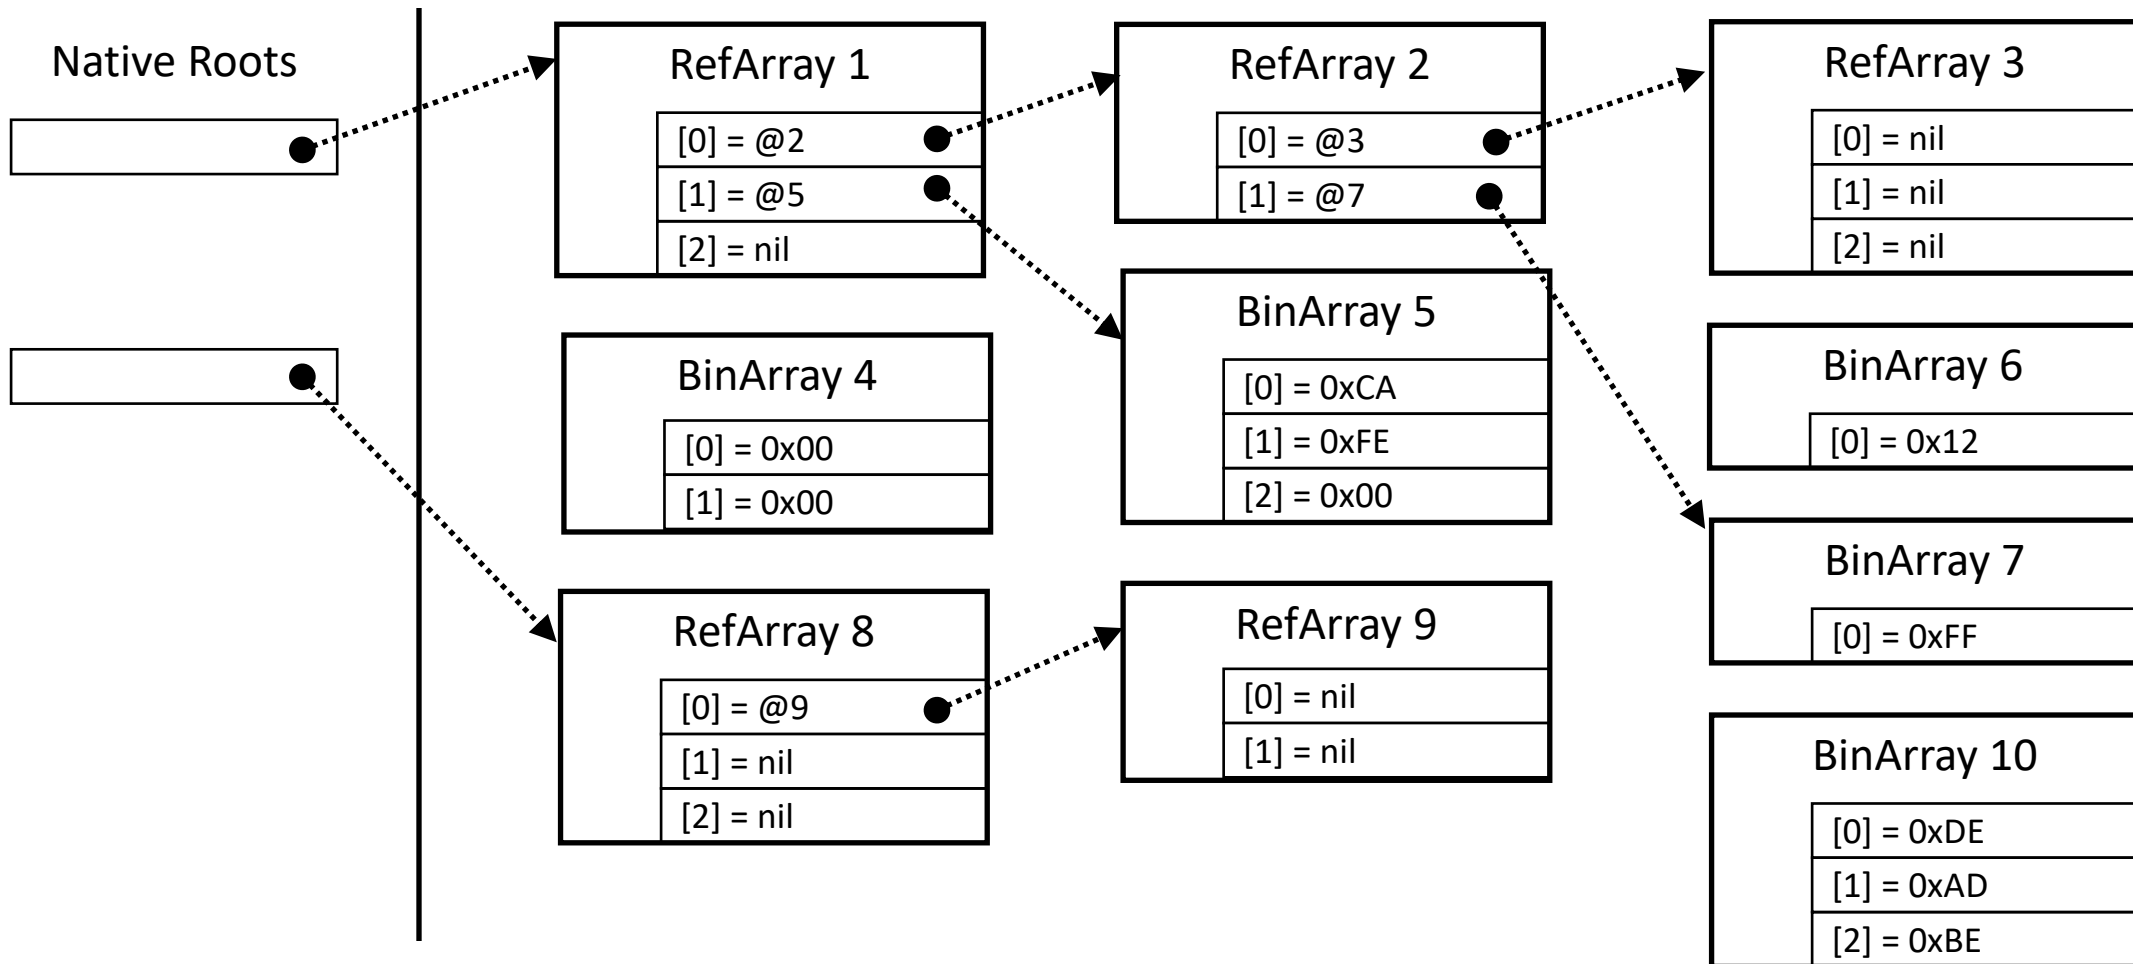
Memory: Graphs of objects



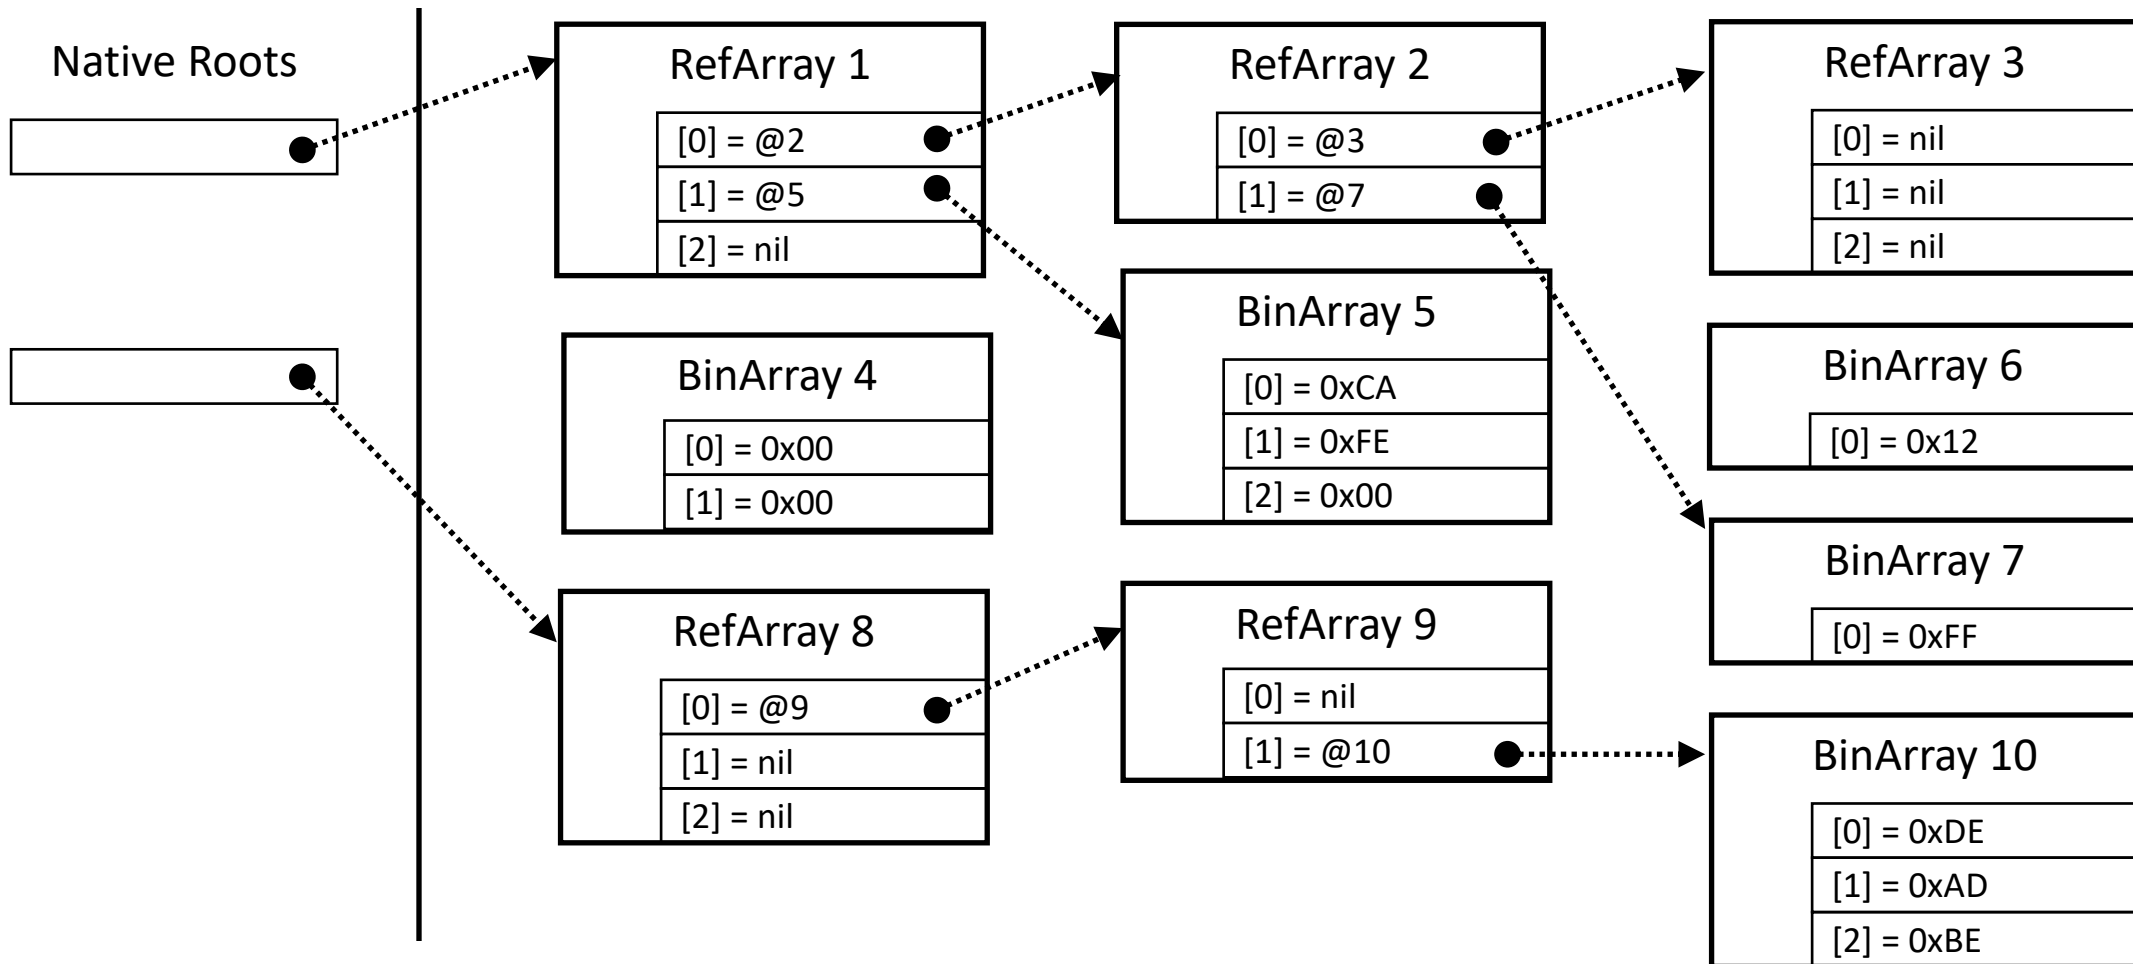
Memory: Graphs of objects



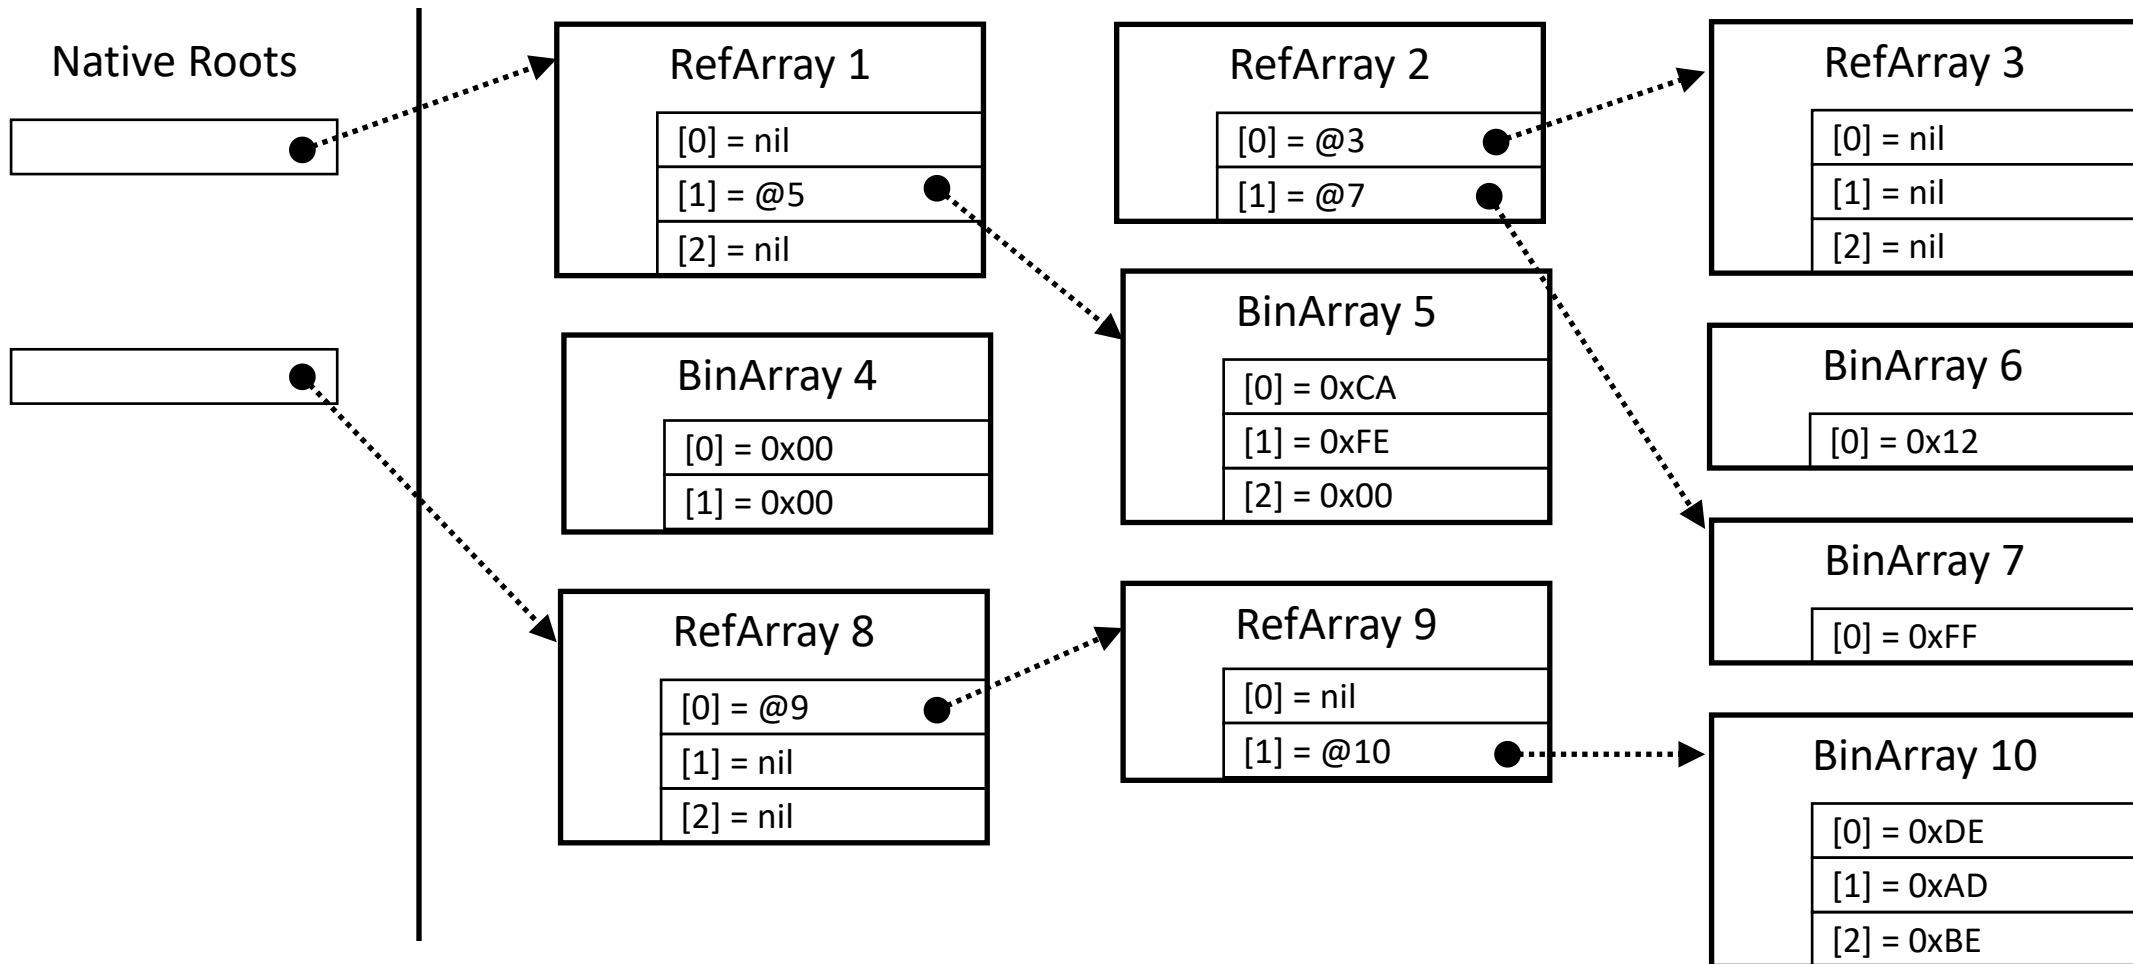
Memory: Graphs of objects



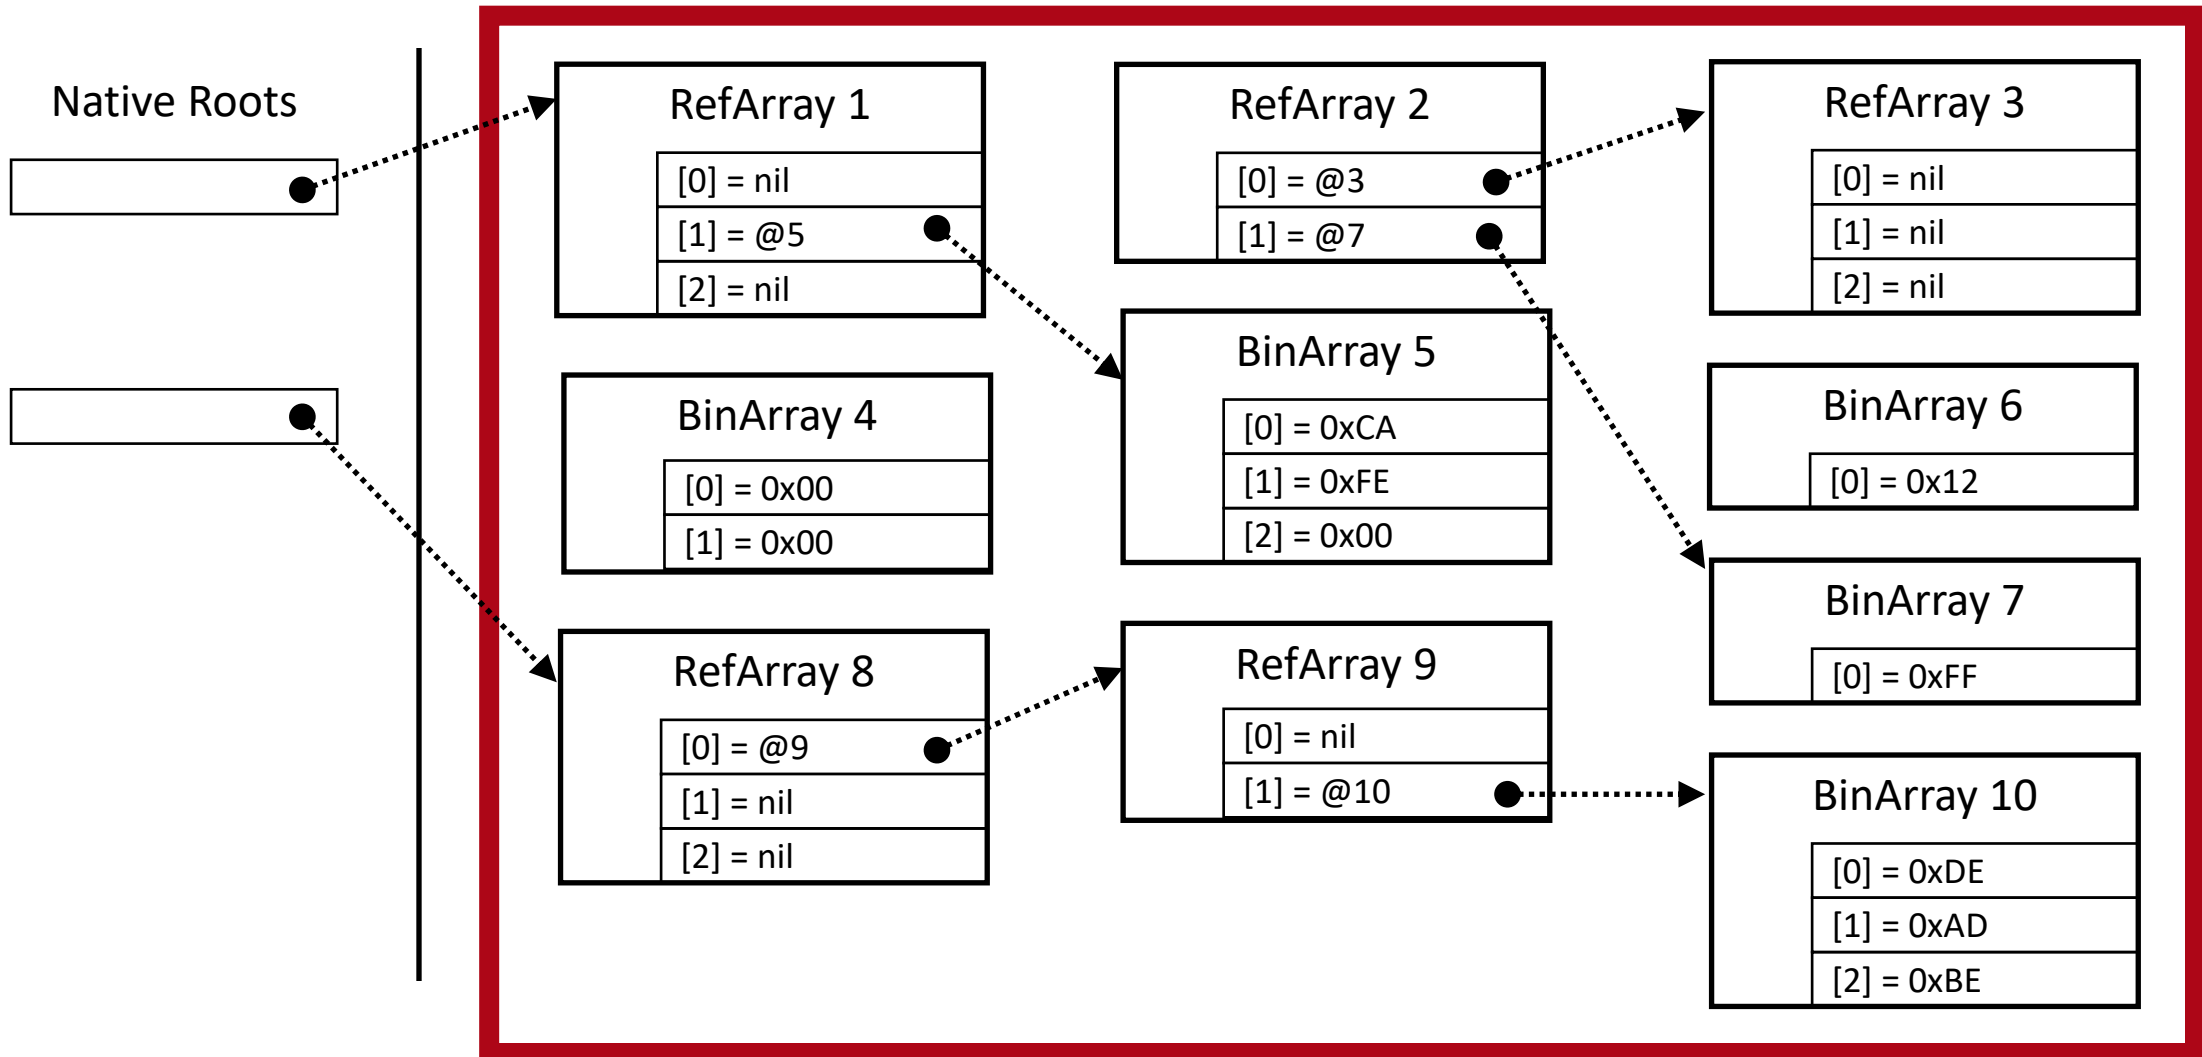
Memory: Graphs of objects



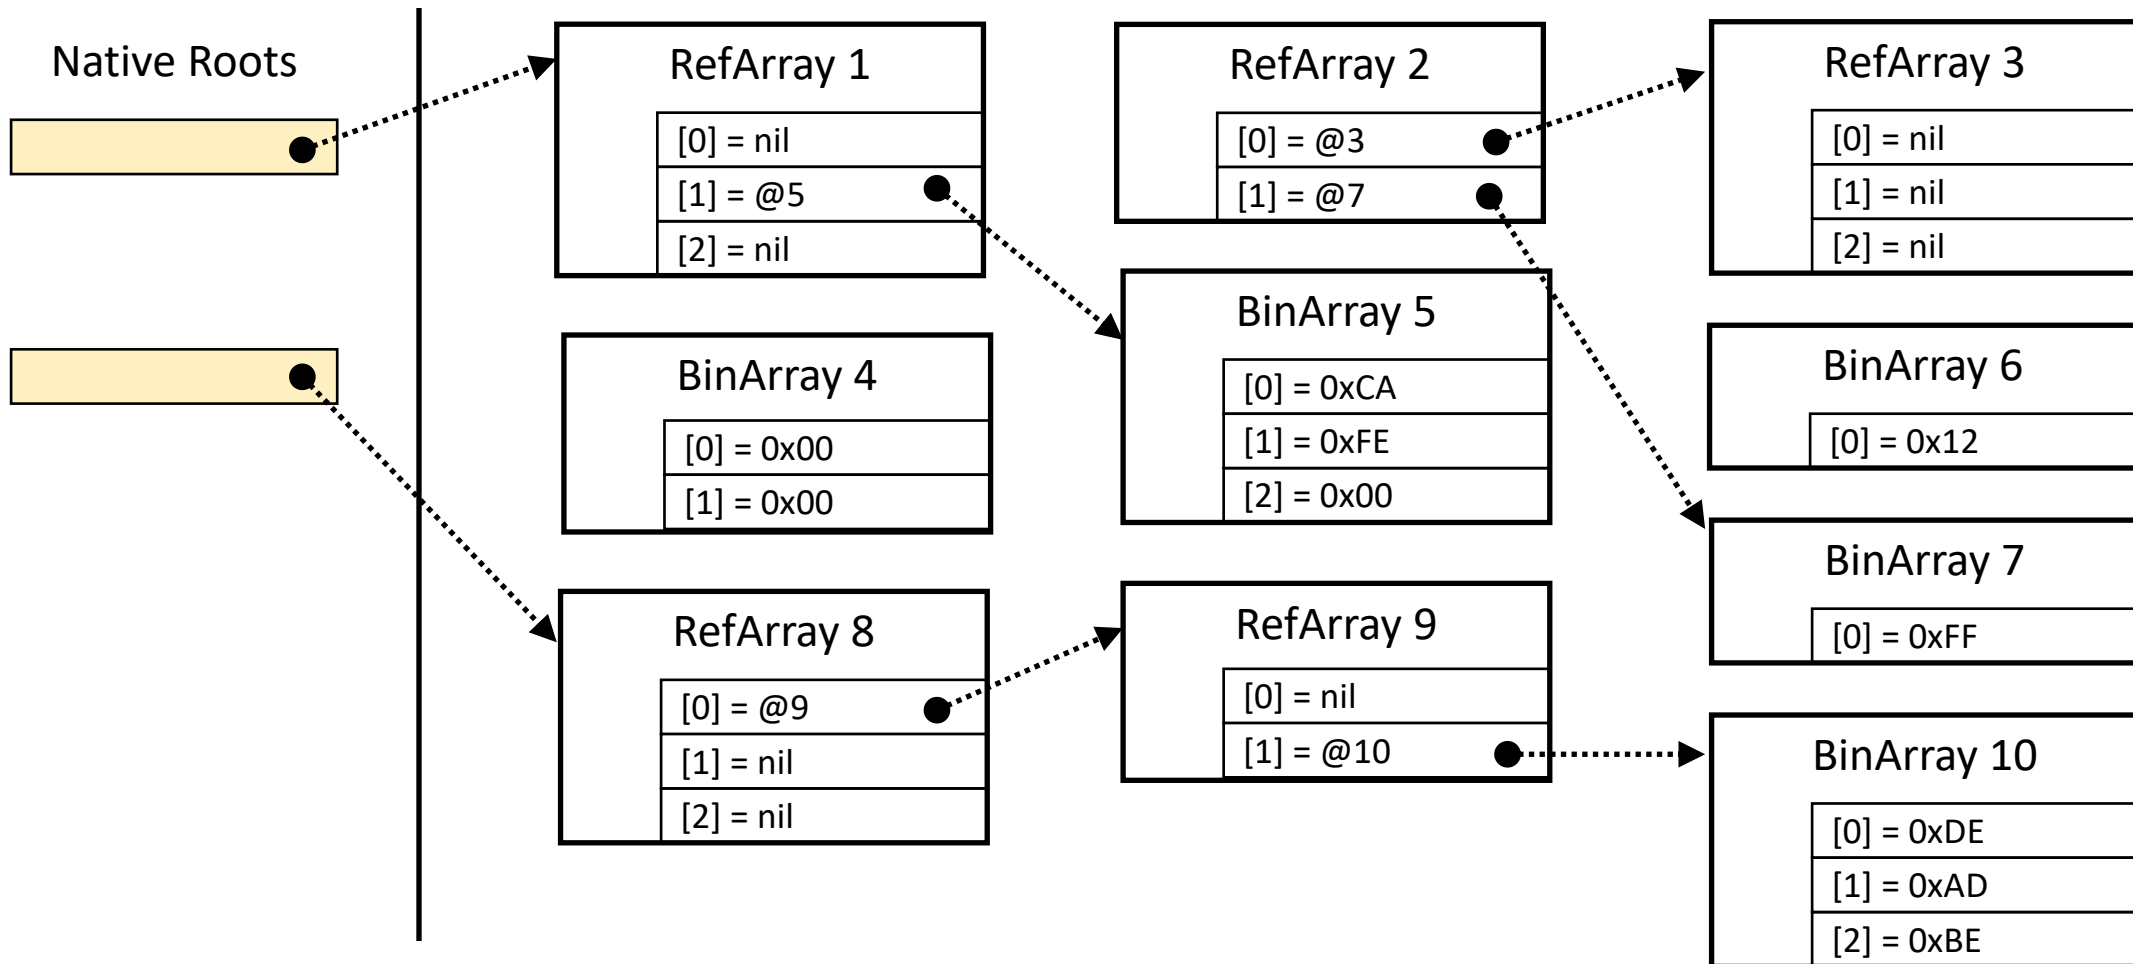
Memory: Graphs of objects



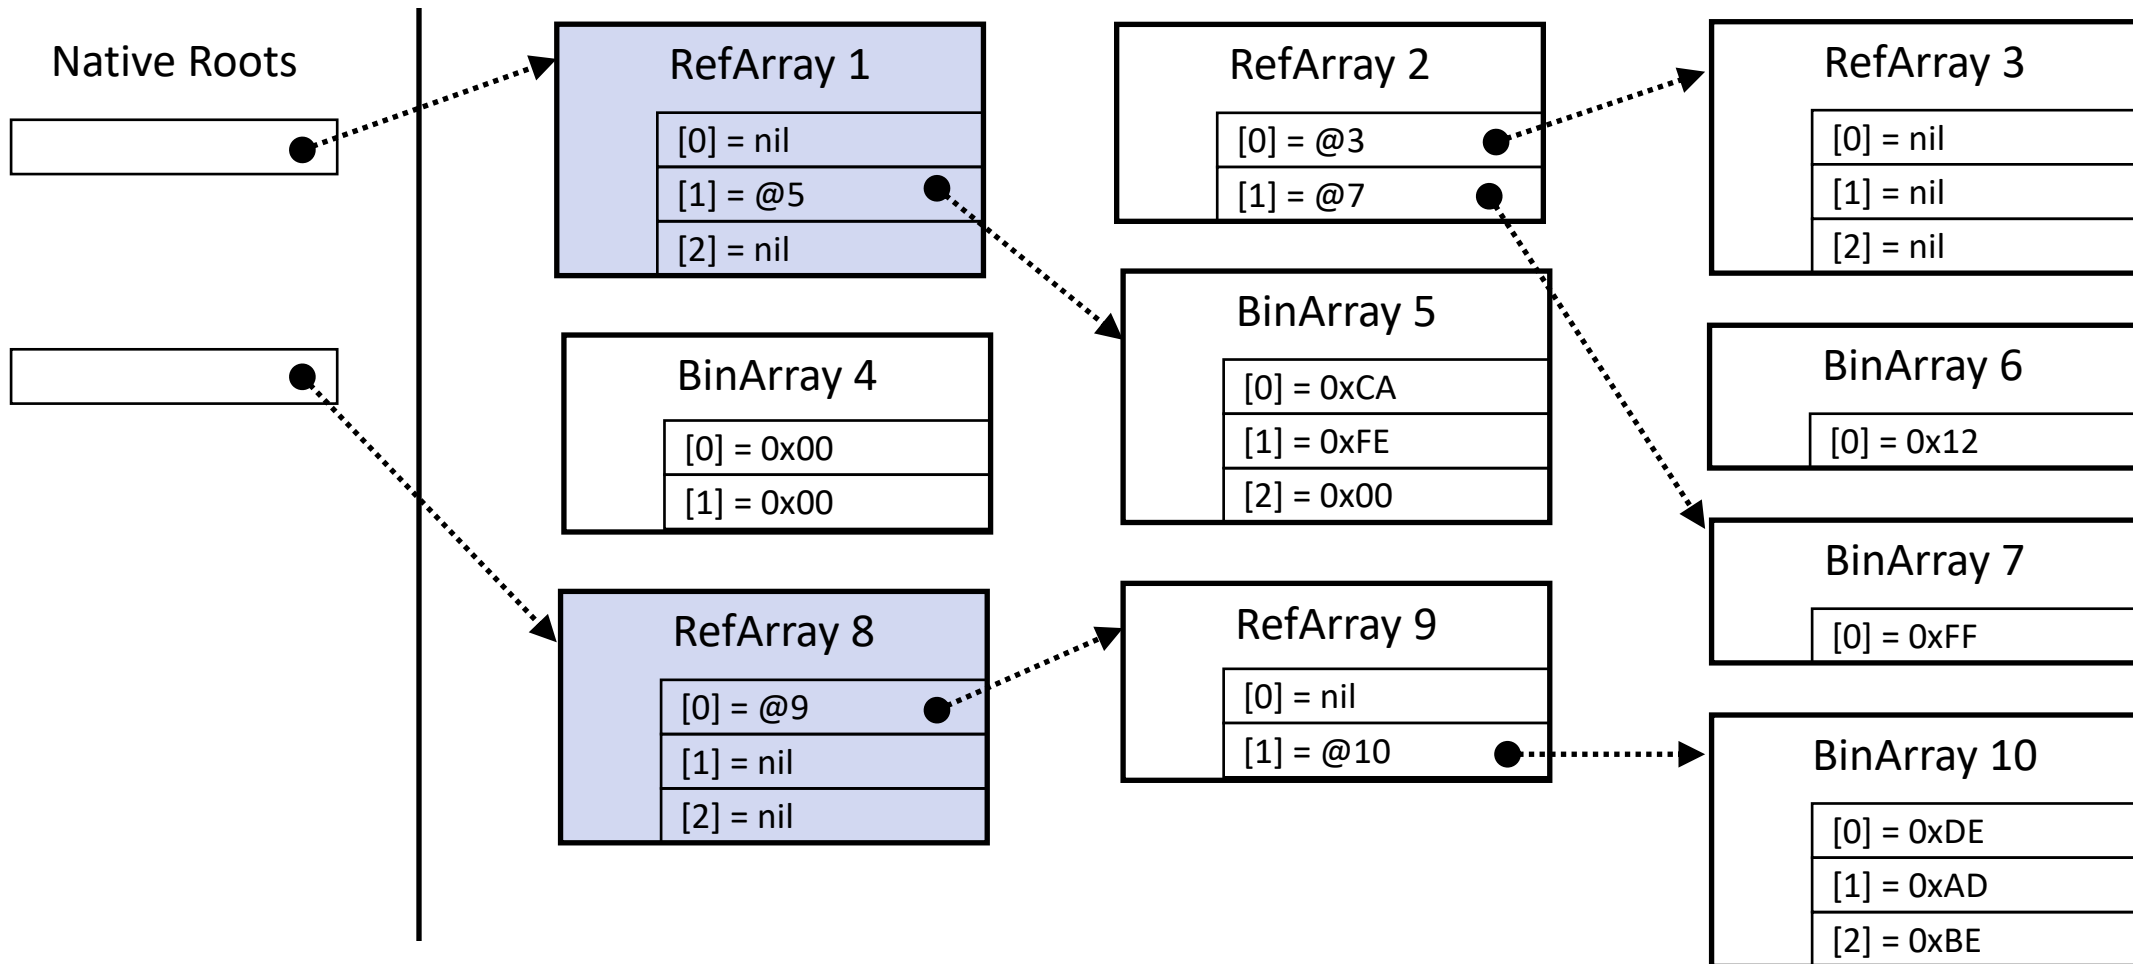
Out of memory!



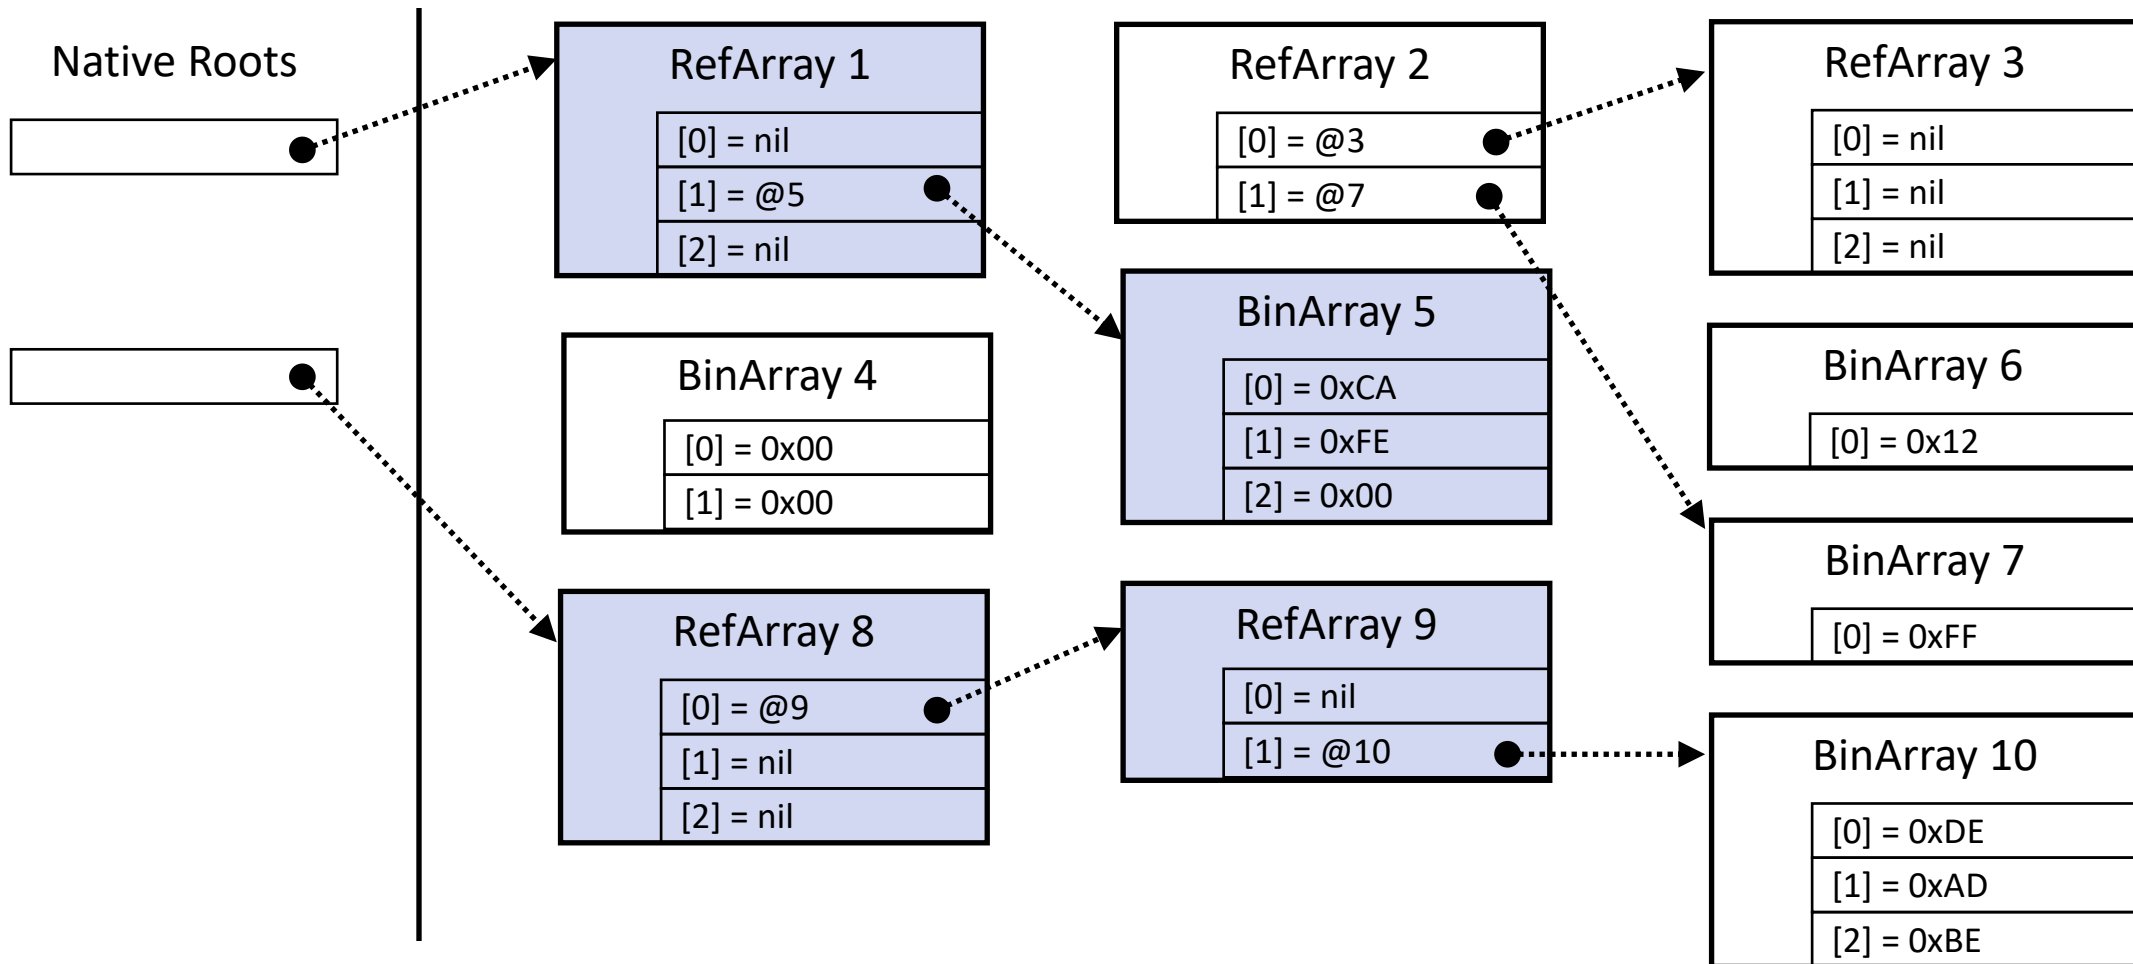
Memory: Graphs of objects



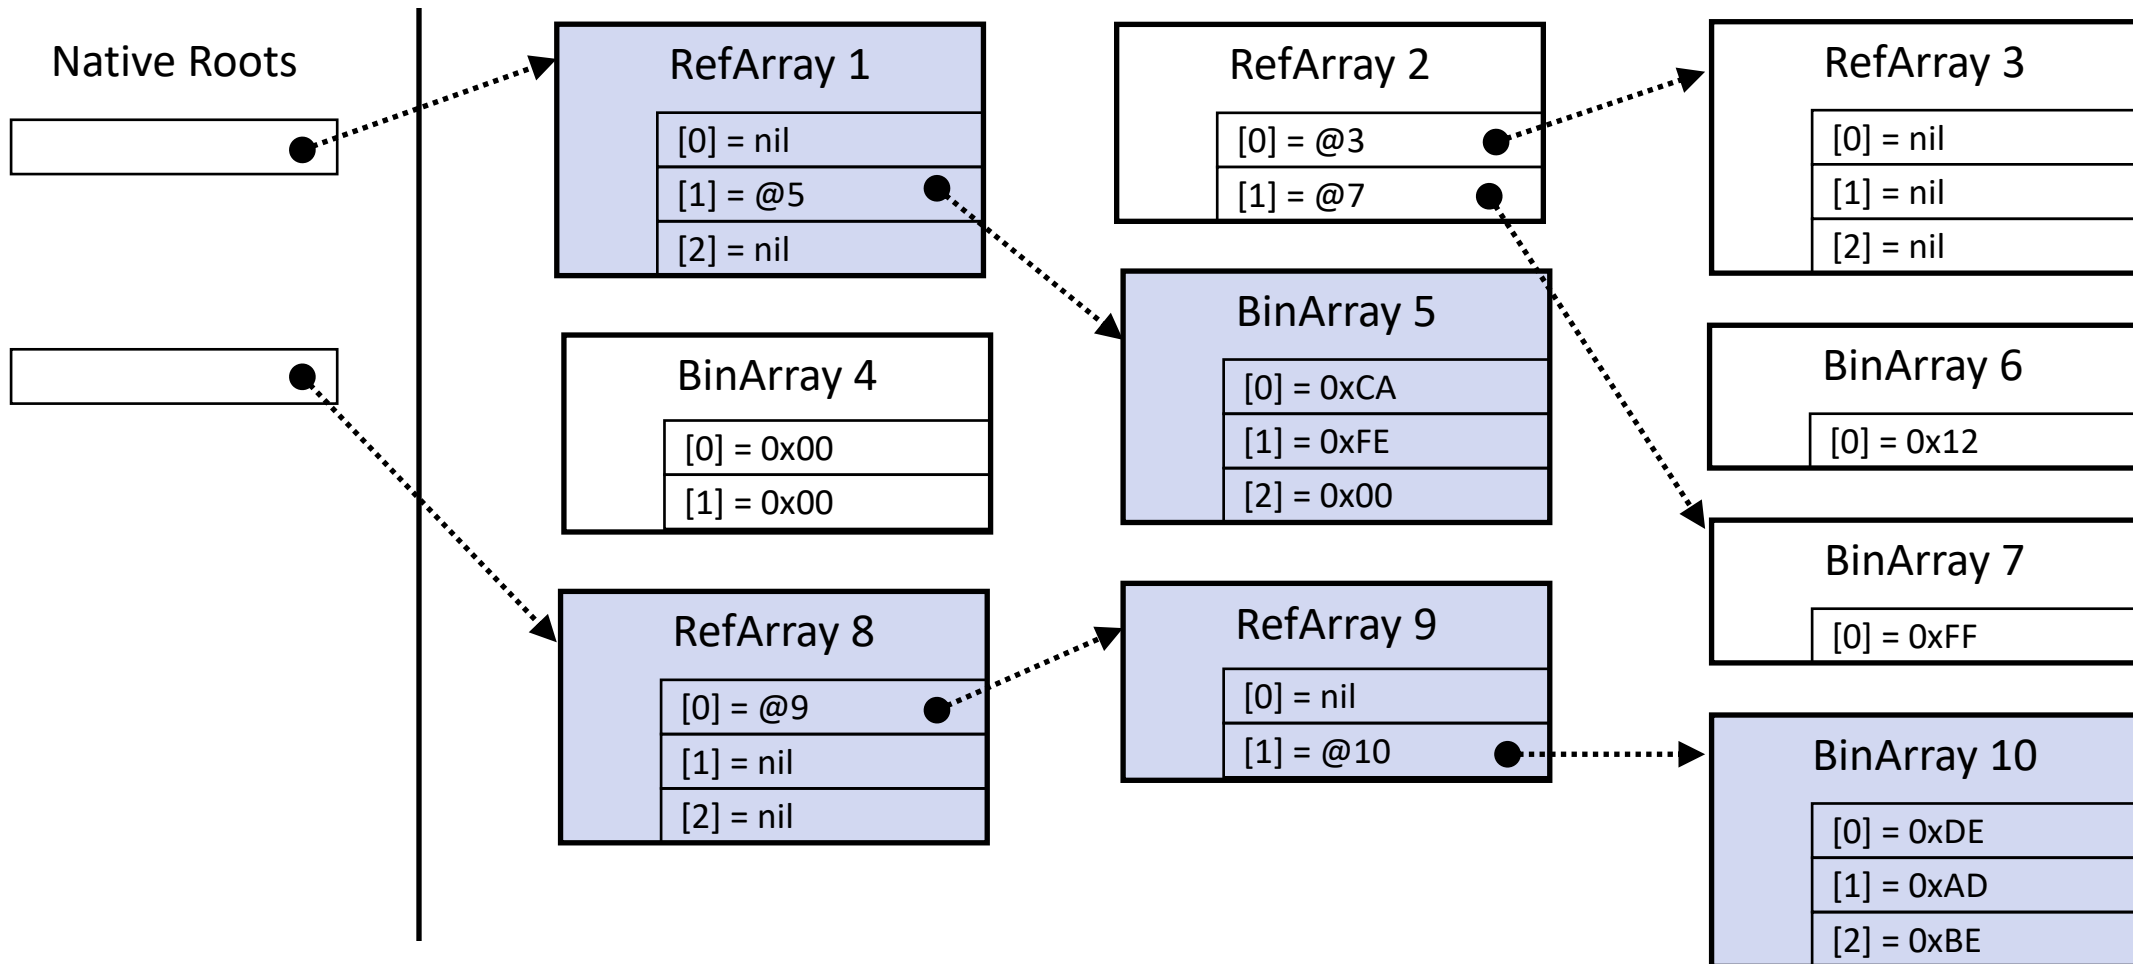
Memory: Graphs of objects



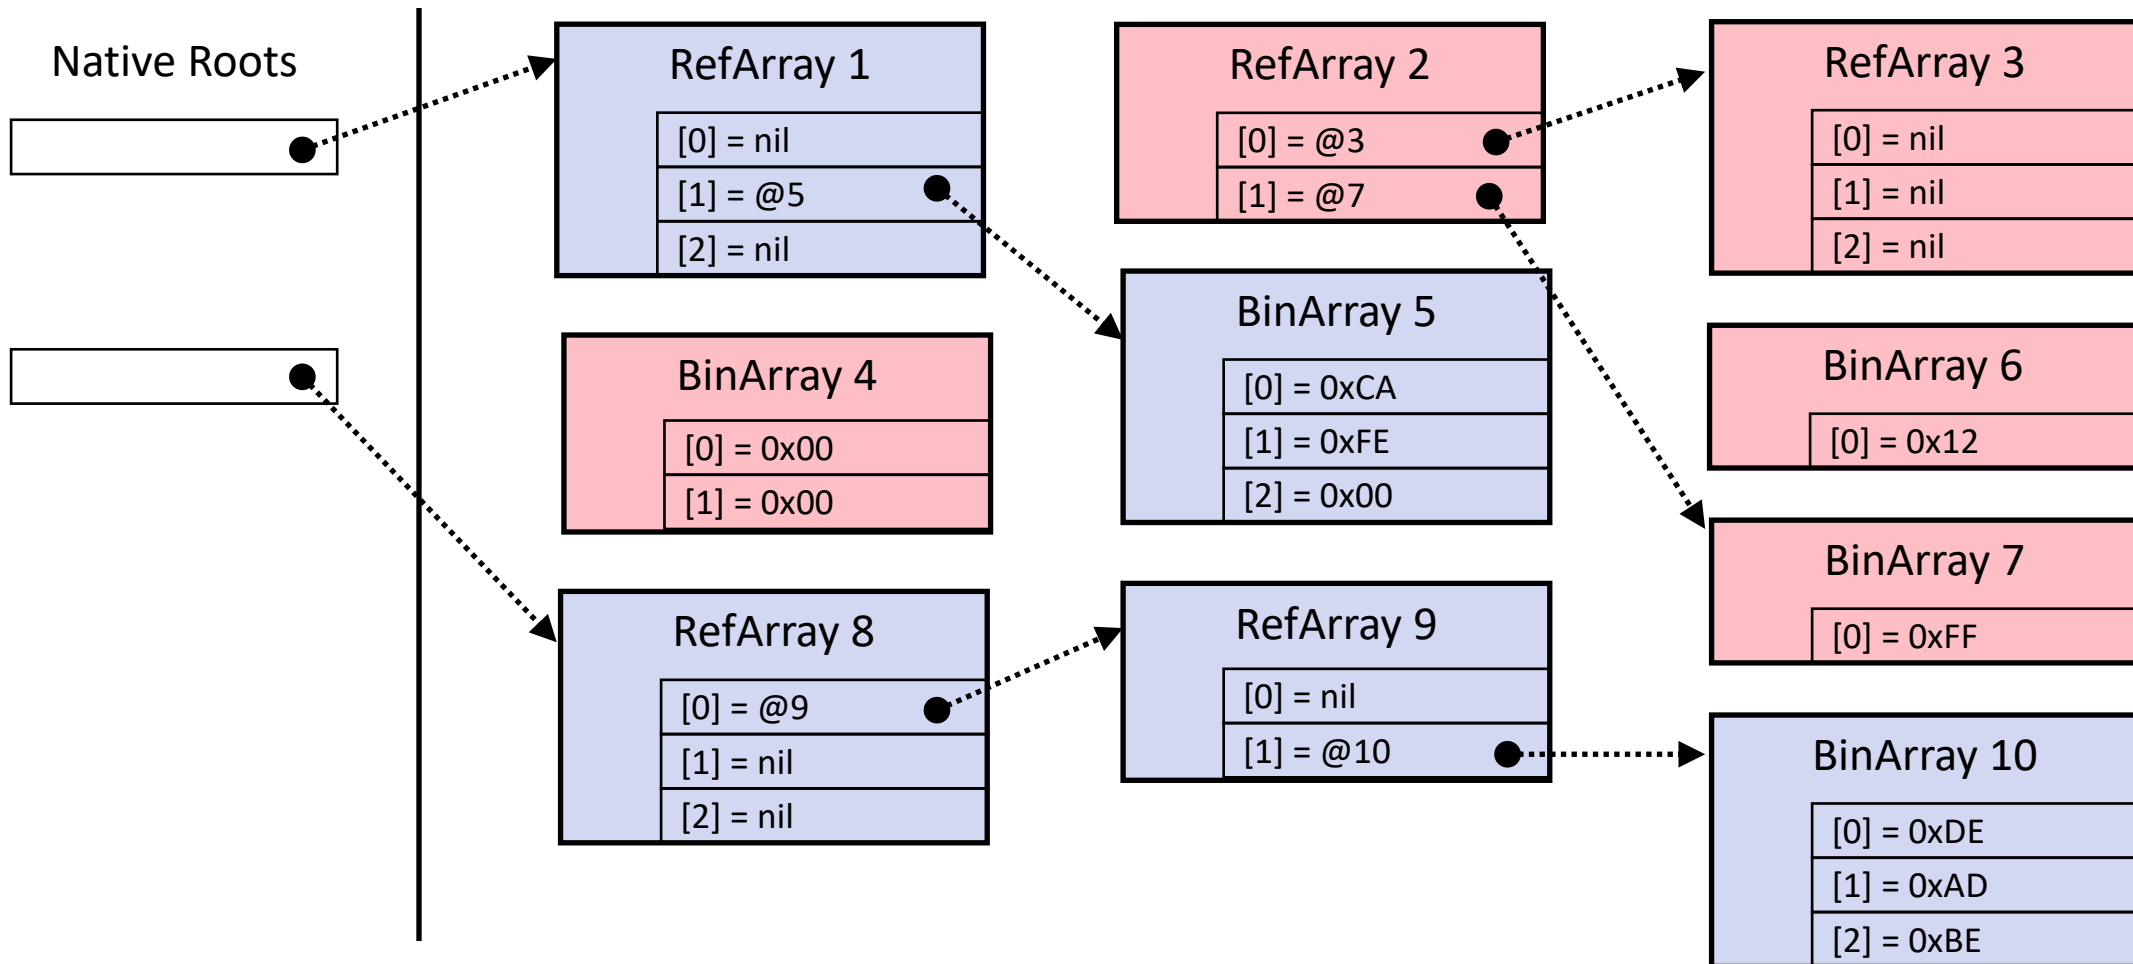
Memory: Graphs of objects



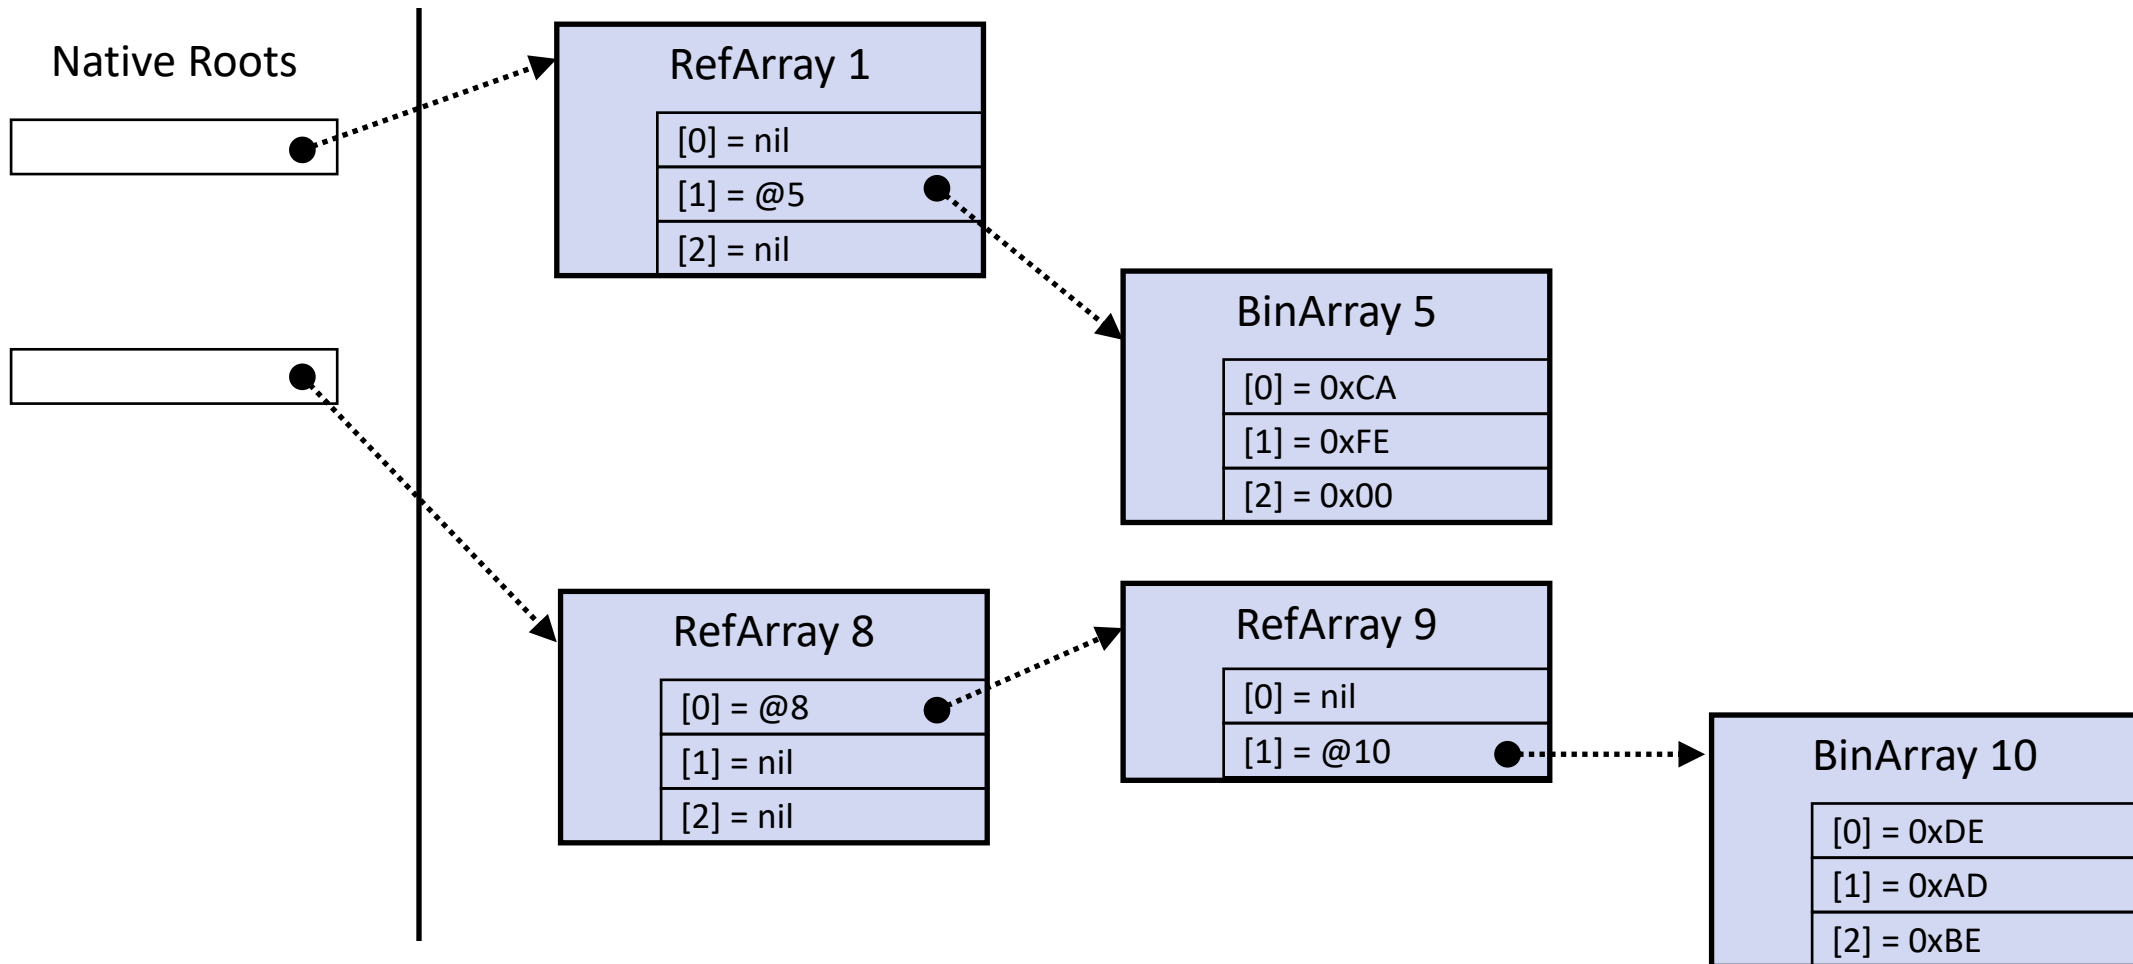
Memory: Graphs of objects



Memory: Graphs of objects



Memory: Graphs of objects



What's going on?

1. The collector is using a classic “mark and sweep” algorithm
 1. Scan roots, mark reachable objects, put them on a work stack
 2. Scan the objects on the works stack, to find new live objects
 3. Build a map of the used and unused portions of the heap
 4. Add unused portions of the heap to the free list
- Memory is reclaimed in bulk, on demand
- Free memory is found in the space “between” live objects
- The GC has no “per-object” free operation (no destructors)

What does the GC need to know?

- Root scanning – what objects are we working with?
- Object size
- How to find references between objects:
 - Slot Location: object + offset
 - Slot Encoding: need to read and write references
- When is the graph changed?

Configuring the GC

- OMR is massively configurable at compile time
- You must teach the GC about your objects and runtime
- Users (that's us!) implement “client” code
- A set of APIs defined by OMR, but implemented by consumers
- Client code is compiled and inlined into OMR
- Clients can incrementally develop their client code to enable new technology

The OMR GC API

An experimental set of APIs for the collector

Initializing the collector

- `OMR::Runtime`
 - Process wide singleton
 - Responsible for initializing the port & thread library
 - Required to bring up the GC subsystem
- `OMR::GC::System`
 - A complete garbage collected heap
 - Static configuration is optionally passed in to the constructor
 - You can bring up multiple heaps per process (hopefully, haha !)
- `OMR::GC::Context`
 - A per-thread GC context, required for most public APIs
 - Provides local heap caches, heap access locks, and rooting utilities

Collector initialization

```
#include <OMR/GC/System.hpp>

// Process-wide singleton
OMR::Runtime runtime;

// Each system contains a unique heap. Per-VM.
OMR::GC::System system(runtime, config);

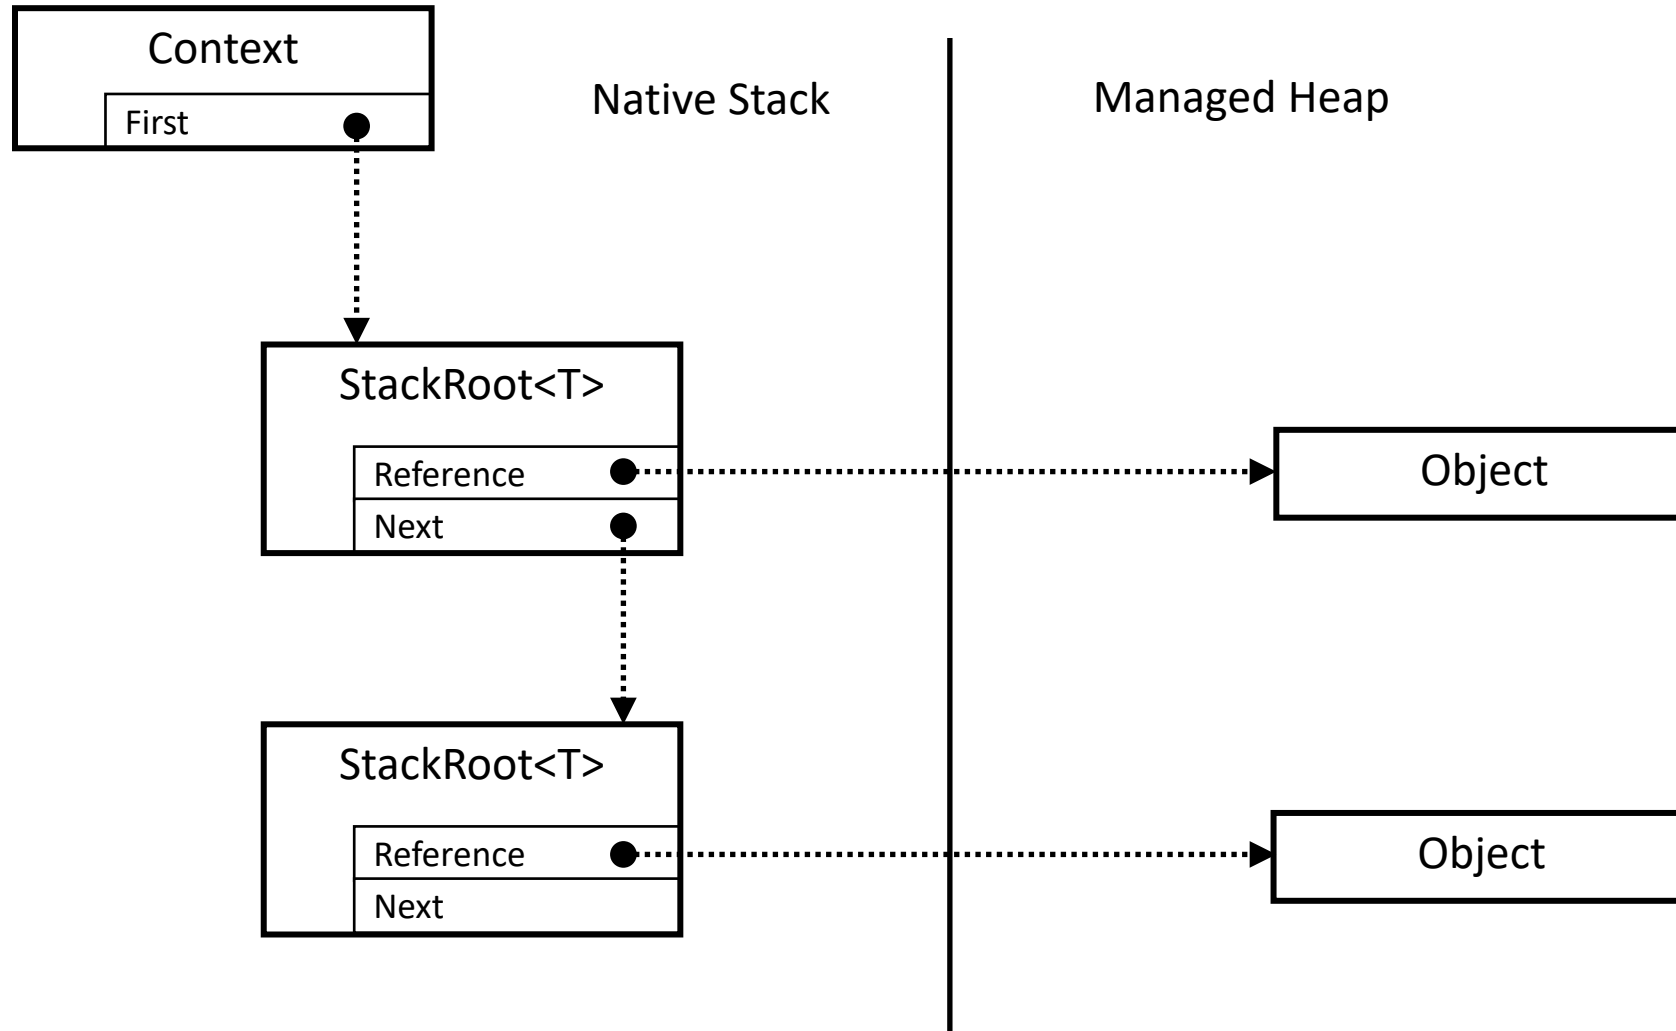
// Thread-local context to the GC::System.
GC::Context cx(sys);
```

The object-oriented allocator

```
template <typename T, typename Init>  
T* OMR::GC::allocate<T>(cx, size, Init init);
```

- The return value is an unsafe heap reference.
- The initializer must put the new allocation into a scannable state
 - IE: set the objects size, and clear any reference slots.
 - The initializer cannot allocate.
- Collections can happen at allocation sites.
 - Do not hold raw heap reference across allocations sites!
 - Use the NoCollect API when it's not safe to collect.

StackRoots: Automatically rooted pointers



StackRoot: Automatically Rooted References

Attached to a specific context, and null by default:

```
StackRoot<Object> root(cx);
```

Assignable, and comparable:

```
StackRoot<Object> r0(cx, allocateObject());  
StackRoot<Object> r1 = r0;  
r0 == r1; // true
```

Have a pointer-like API:

```
Root->field = 42;  
(*root).field = 42;
```

Stack Roots have LIFO semantics and must be allocated on the stack.

Finding slots in Objects

- We need to show the GC how and where GC refs are stored
- We implement an object scanner that can notify GC visitors about edges between objects
- We give the visitor slot handles (pointers to slots).
- The GC uses these handles to read/write references from object slots.
- The `OMR::GC::RefSlotHandle` can be used for slots containing plain, untagged, full-width addresses
- Clients can provide their own slot handle types for defining custom read/write operations.

Let's get started!

Coffee break!

Heap Compaction


Heap Compaction

- Over time, heap memory becomes fragmented
- Fragmentation is bad
- The collector can slide left compaction
- Eliminates heap fragmentation
- Extremely important for long lived applications

The Heap

 Free Space


 Objects




The Heap

 Free Space


 Objects



The Heap

 Free Space


 Objects



The Heap

 Free Space


 Objects



The Heap

 Free Space

 Objects



Generational collection
(Scavenging objects)

What is generational GC?

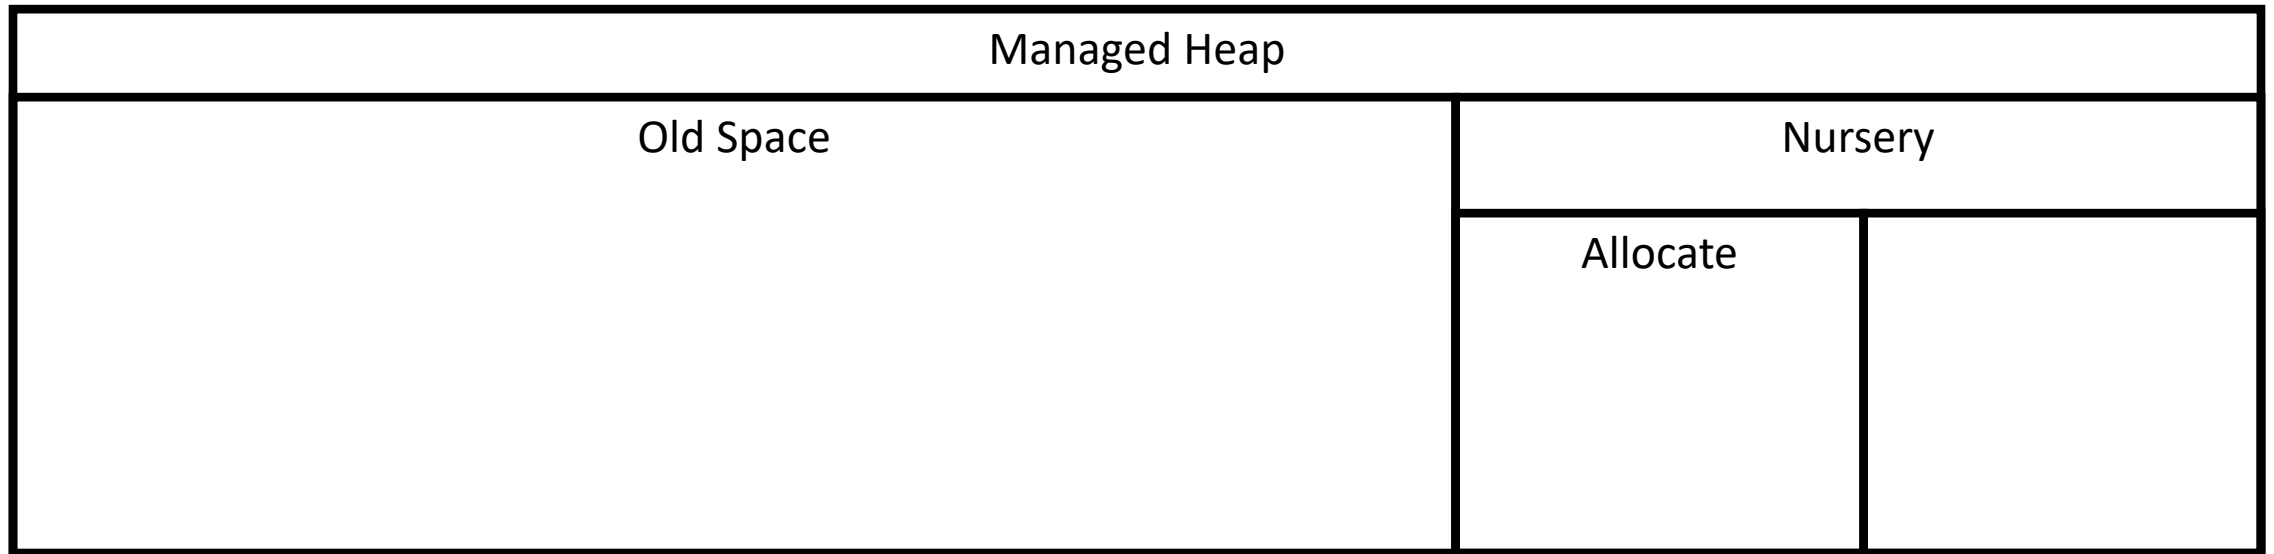
- Weak generational hypothesis:
 - Young objects are more likely to die, or
 - The longer an object lives, the more likely it is to survive.
- The plan: scan only newly allocated objects
- Old objects will survive
- When objects survive long enough, tenure to old-generation
- Also known as “local collection”

Remembering old objects

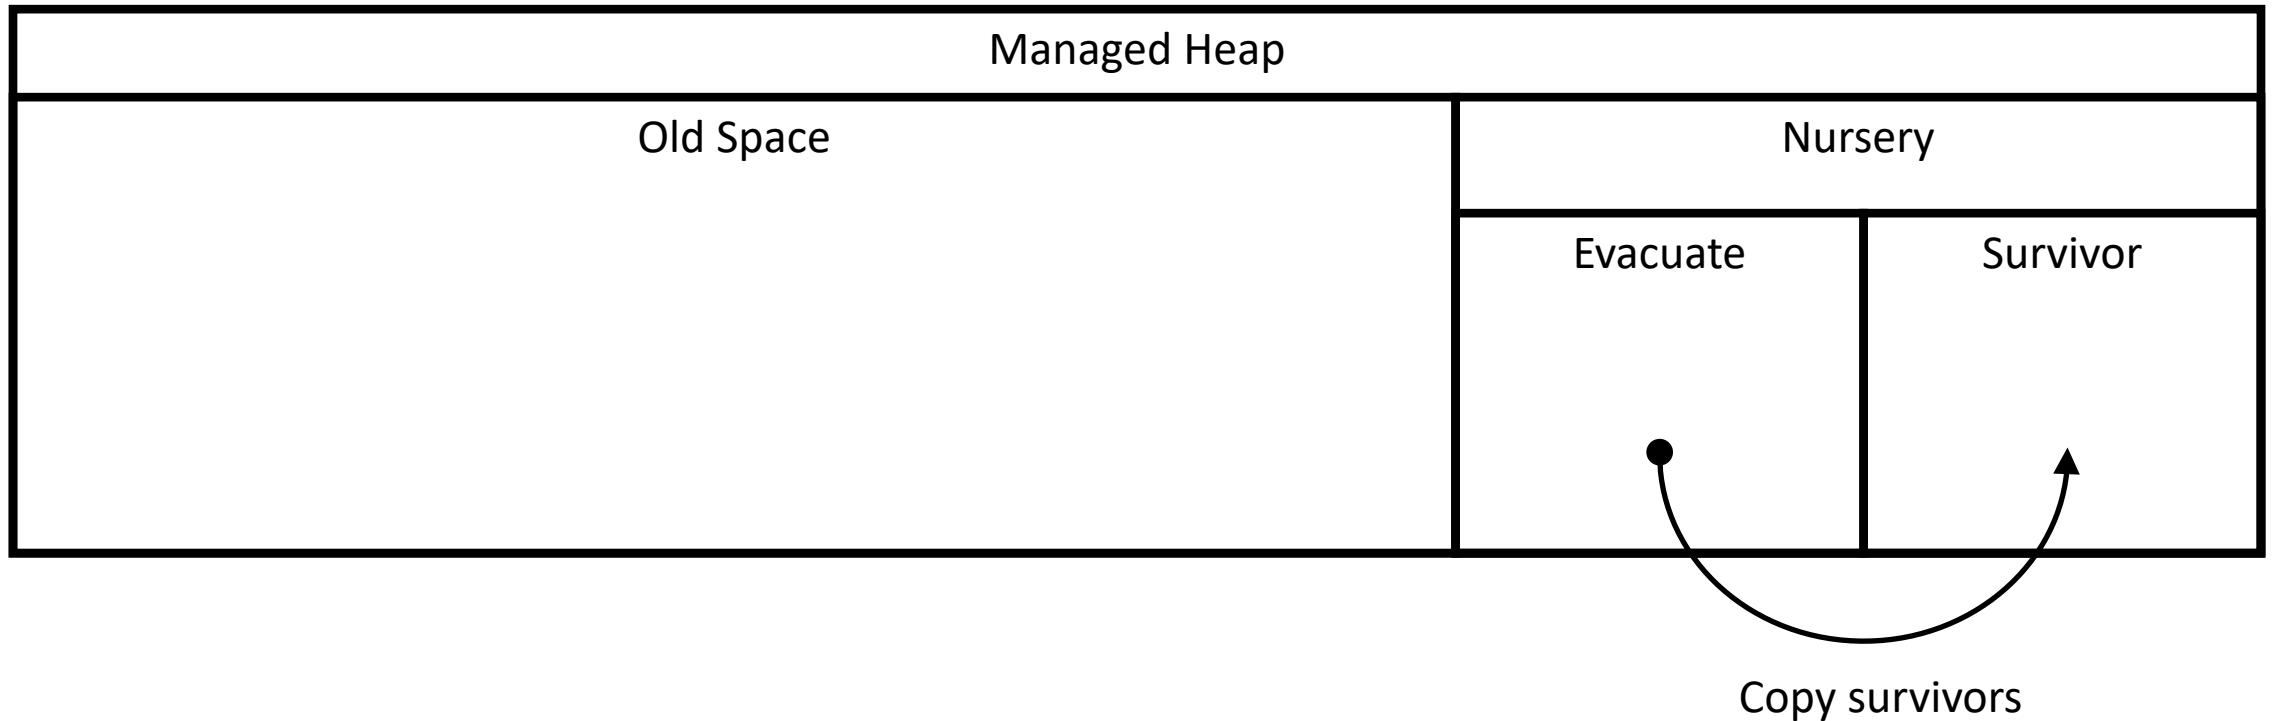
- When an old object references a new object, we “remember” it
- During a local collection old objects will survive
- In a local collection, treat remembered objects as roots
- How do we find old -> new references?
- Use a “write barrier” to track all object graph updates
- Every time a reference is stored:
 - Is the referrer in old space?
 - Is the referent in new space?

=> Remember the referrer

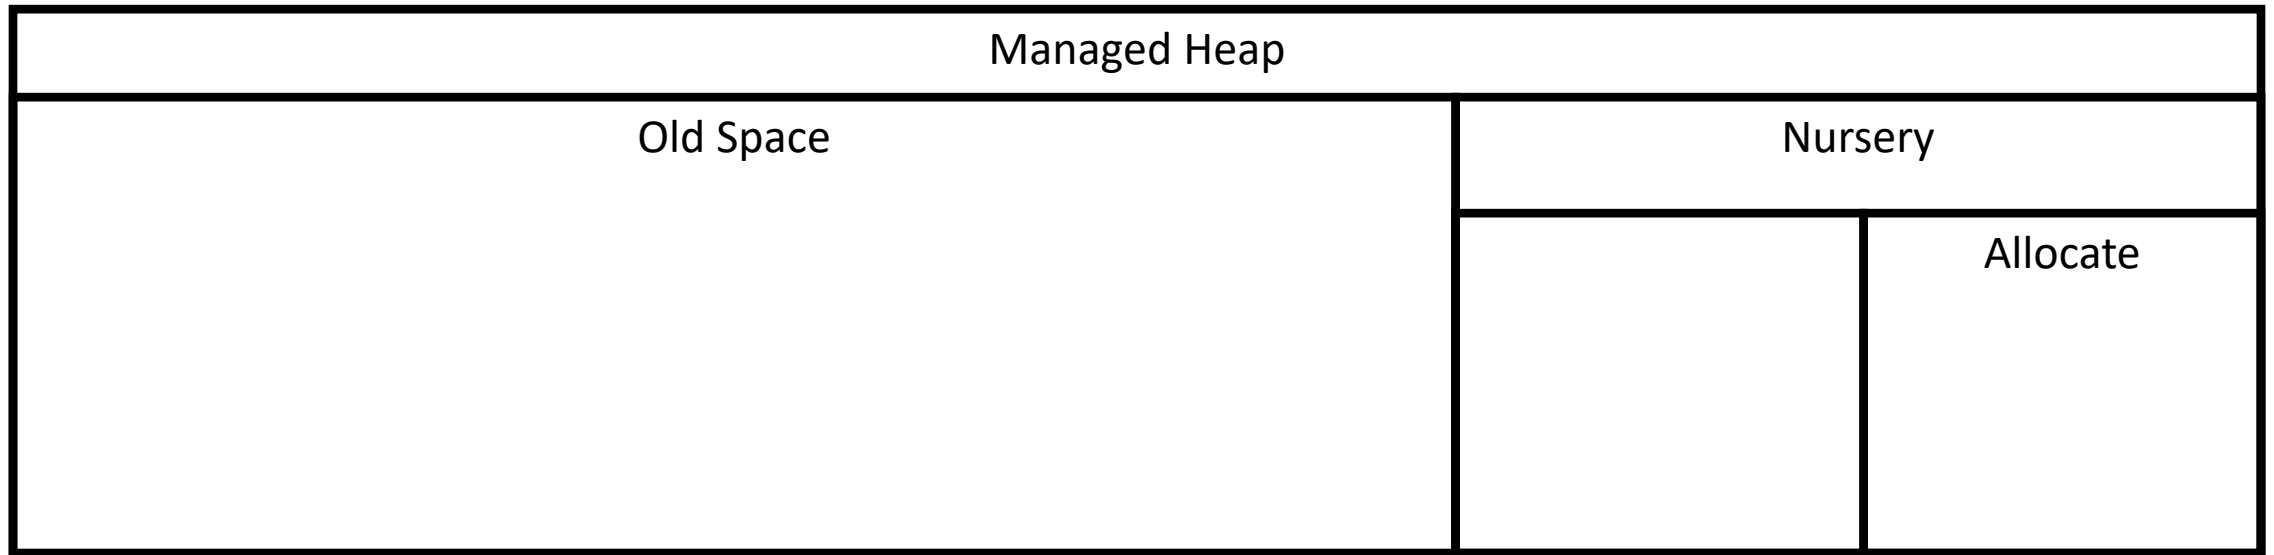
The Generational Heap



The Generational Heap: Scavenge



The Generational Heap



OK, Back to work

Thank You!!